Center for Spatially Integrated Social Science

Spatial Regression Analysis in R
A Workbook

Luc Anselin

Spatial Analysis Laboratory
Department of Geography
University of Illinois, Urbana-Champaign
Urbana, IL 61801

http://sal.uiuc.edu/

Center for Spatially Integrated Social Science

http://www.csiss.org/

Revised Version, May 10, 2007

Contents

List of Figures

Preface

This workbook contains a set of laboratory exercises initally developed for graduate courses in *Spatial Econometrics* and *Advanced Spatial Analysis* at the University of Illinois, as well as for the ICPSR Summer Program course on *Spatial Regression Analysis.* It consists of a series of brief tutorials and worked examples using R and its packages spdep for spatial regression analysis and spgwr for geographically weighted regression.

Some of these materials were included in earlier tutorials available on the SAL web site. In addition, the workbook incorporates all exercise materials for the course ACE 492SE, *Spatial Econometrics*, offered during the Fall 2003 and Spring 2005 semesters in the Department of Agricultural and Consumer Economics at the University of Illinois, Urbana-Champaign. There may be slight discrepancies due to changes in the version of R and spdep. In case of doubt, the latest document available on the SAL web site should always be referred to as it supersedes all previous tutorial materials.

The examples and practice exercises use the sample data sets that are available from the SAL "stuff" web site. They are listed on and can be downloaded from http://sal.uiuc.edu/stuff/stuff-sum/data. The main purpose of these sample data is to illustrate the features of the software. Readers are strongly encouraged to use their own data sets for the practice exercises.

Warning: These notes are *not* polished, and remain very much in bullet style. They are intended to provide hints and additional pointers beyond what is available in the help files and the other tutorials. They also serve to define a sequence of tasks and techniques that can be covered during a given lab session. They are "in progress" and are still undergoing changes.

The development of this workbook has been facilitated by the continued research support through the U.S. National Science Foundation grant BCS-9978058 to the Center for Spatially Integrated Social Science (CSISS).

Exercise 1

Getting Started in R

1.1 Objectives

These notes are the first in a series to help you get started and working in R to carry out spatial regression analysis. The objective is that you work through the various examples at your own pace. This chapter and later ones will refer to the two short R tutorials available on SAL "stuff" web site, as well as to the *Introduction to R* (Venables et al. 2004), available from the CRAN site. I will refer to Venables et al. (2004) as the "Introduction" tutorial. I will refer to the other tutorials as the "Data" tutorial (*Data and spatial weights in spdep*, Anselin 2003a), and the "Regression" tutorial (*An introduction to spatial regression analysis in R,* Anselin 2003b). Further details on the spdep package can be found in Bivand and Gebhardt (2000), Bivand (2001, 2002a,b), and Bivand and Portnov (2004).

1.2 Conventions and General Advice

Some notational conventions:

- The command line will be shown as > followed by an R command.

- Command line contents will be in `monospaced` font.

- File names and file path names will be in *italics*.

A few general pointers:

- Create a separate directory for your work in R. Specify that directory as your starting directory in the R shortcut, as illustrated on p. 2 of the *Regression* tutorial.

- Know how to quit: `>q()` from the command line, or `File > Exit` using the R GUI.

- Know how to get help. One option is to invoke `>help.start()`. This launches your default browser as an interface to HTML help pages. Another option is to invoke help on a specific function or data set as `>help(name of function)`. A powerful feature is to use fuzzy search, so that you don't have to know the exact command you need (if that is what you are looking for in the first place).

- A useful way to find out more how a command works is to cut and past sample code that is contained in the help.

- The assignment symbol is `<-`. This assigns the result of what it done on the RHS to the "object" on the LHS. Often it may seem like nothing happens after you run a function. Typically that is because the result was assigned to an object without an explicit call to `print` or `summary`. After this assignment you manipulate the object, e.g., to print or summarize the results.

1.3   Sample Session

As a first orientation, work through the sample sesssion in Venables et al. (2004), Appendix A. This is a good way to get an idea of what R can do and how it works.

You can either follow along in the pdf file, or you can open the Introduction in HTML format. To do the latter, invoke the HTML help command by selecting (`>help.start()`), and click on the entry "An Introduction to R," and select "A sample session".

This set of examples requires a specific sample data set, *morley.tab*. The tutorial asks you to copy that file to your working directory. In the current version of R (2.0.1) all sample data sets are contained in the `datasets` package. To make it available, enter `>library(datasets)` followed by `>data(morley)`. All the variables in the morley data sets should then be available. Once you complete this, either type the commands from the sample session on the command line (to get used to the experience of unforgiven typos) or cut and paste from the HTML file. Needless to say, the latter is the preferred approach. Note that when using the morley data set from the `datasets` package, you need to `attach(morley)` to get direct access to the variables. At that point, you no longer need to specify `data=mm` as shown

in the sample code (if you copy the commands from the html help file, this will give an error message).

After going through the sample session, you may want to experiment with some of the commands. Look up the specifics using `help( )` and cut and paste code from the examples to see how it works. You may also want to run some `demo( )` for applications (not all of them have a `demo` function implemented). For example, a nice illustration of the graphics capability is `demo(graphics)`.

Exercise 2

Getting Your Data and Weights Into R

2.1 Objectives

The purpose of this exercise is to provide guidance in turning your data and spatial weights files into usable objects in R, and to carry out some basic data and weights manipulations. The assumption is that you have the data originally as shape files and will construct the weights files in *GeoDa*. Of course, R has its own set of functions to create an analyze spatial weights, and we will consider those in a separate exercise (see also the *Data* tutorial).

You should create a new working directory for your current R project and make sure all files (data files and weights files) are copied to that directory. This is not absolutely necessary, but it's a good idea, since it avoids problems with path names, etc.

You will be using the POLICE sample data set from the SAL data archive. Make sure you download the file and unzip it into your working directory. After unzipping the file, check the `police.html` file to become familiar with the contents of the data set.

2.2 Getting Data and Weights File Ready

You first need to export the data from the shape file format to an ascii file that R can read. You also need to create some weights files using *GeoDa*. A similar exercise is described in the *Data* tutorial, using older versions of R and `spdep`.

2.2.1 Exporting Data with GeoDa

Load the `police.shp` file into *GeoDa*, with `FIPSNO` as the `Key Variable`. Select `Tools > Data Export > Ascii`. In the dialog, specify `police.dbf` as the input file and `police.txt` as the output file. Select all variables, except `AREA`, `CNTY_`, `CNTY_ID` and `PERIMETER`. Click on `Export` and `Done` to finish. Check the contents of the `police.txt` file and make sure it shows `82,12` on the first line and the variable names on the second line. The data, record by record, and comma-delimited make up the remainder. Copy the file to the R working directory if it is not there already.

2.2.2 Creating the Spatial Weights in GeoDa

For this exercise, create two spatial weights files for the POLICE data set, one based on rook contiguity, the other on 5 nearest neighbors. Use `FIPNO` as the `ID` variable. Call these files `policerook.gal` and `policek5.gwt`, respectively. See the instructions in the *GeoDa* Workbook for details. Make sure you save (or copy) the files to the working directory for your R project.

2.3 Creating a Data Frame

2.3.1 Command Line

In R, the easiest way to create a data frame from a comma-delimited ascii file is to use a variant of the `read.table` function that takes into account the presence of the comma. This is the `read.csv` function. Check out `>help(read.csv)` to see all the details and options. The files created by *GeoDa* contain a header line (`header=TRUE`) and an extra first line that R doesn't need (the line with the `82,12`). You avoid reading this line by setting the `skip` option to 1. The proper command to read the file is then:

```
>police <- read.csv("police.txt",header=TRUE,skip=1)
```

As usual, nothing seems to happen when you enter this, but you can check the presence of the new `police` object with a `summary` command:

```
> summary(police)
     FIPSNO          POLICE            POP              TAX
 Min.   :28001   Min.   :  49.0   Min.   : 2500   Min.   : 81.0
 1st Qu.:28042   1st Qu.: 291.5   1st Qu.: 13275   1st Qu.:106.2
 Median :28082   Median : 432.5   Median : 20700   Median :125.5
```

```
Mean    :28082   Mean    :  927.8   Mean    : 31311   Mean    :142.2
3rd Qu.:28122   3rd Qu.:  955.8   3rd Qu.: 35850   3rd Qu.:156.5
Max.    :28163   Max.    :10971.0   Max.    :253100   Max.    :674.0
...
```

Also, an `>ls()` command will reveal `"police"` as one of the objects in your work space.

You can specify a variable to serve as the row name in the data frame, but I usually prefer not to do that. The option is `row.names="varname"`.

2.3.2   Writing a Function

The real power of R is obtained through writing functions that carry out repetitive actions. To practice writing functions, you will write a simple procedure to turn a *GeoDa* text output file into an R data frame. You will then "compile" it to make it part of your R environment. At that point, you can invoke it in the same way as any other built-in R function.

Use any text editor (e.g., Notepad or similar) to create a text file, say `read.geoda.R`. Even if you use a "real" word processor, make sure to save the file as *pure* text, so that no extra characters sneak in.

Enter the following lines (lines starting with # are comments and you can skip those to minimize typing):

```
# read.geoda
# helper function to read GeoDa export files
# specify input file = file
# default is no row names, specify row names as second parameter if needed
# example: balt <- read.geoda("baltim.txt")
#                  balt <- read.geoda("baltim.txt","STATION")

read.geoda <- function(file,row.names=NULL)
{
        read.csv(file=file,header=TRUE,skip=1,row.names=row.names)
}
```

The `read.geoda` function simply calls the standard `read.csv` function with some preset options. Note how the `file` is a required parameter, but the `row.names` can be omitted, since it is set to `NULL` as the default. The way the `row.names` is used in the argument list requires only that the name of the variable be specified, no need to write `row.names=`.

To make the function part of your environment, you must first "compile" it with the R `source` command. Make sure the file `read.geoda.R` is in your working directory and enter:

```
> source("read.geoda.R")
```

The `read.geoda` function is now available and can be invoked directly. Note that the R file extension is only for the text file that contains the code, the function as defined in that file does *not* have an R extension. So, don't use `read.geoda.R` as the function name. To practice, first remove the current `police` object using the `>rm(police)` command. Check that it is really gone by trying a `summary(police)`. You should get

```
> summary(police)
Error in summary(police) : Object "police" not found
```

Next, recreate the `police` object with:

```
> police <- read.geoda("police.txt")
```

You can again make sure it is there with a `summary` command or by using `ls()`.

2.4   Creating a Neighbor List from a GAL File

To manipulate spatial weights, you will need to load the `spdep` package. Of course, this assumes that you previously installed it (see the R tutorials on how to do this). To load the package, use the `library` command as illustrated below:

```
> library(spdep)
Loading required package: tripack
Loading required package: sp
Loading required package: maptools
Loading required package: foreign
Loading required package: boot
Loading required package: Matrix
Loading required package: lattice

Attaching package: 'lattice'
```

```
The following object(s) are masked from package:boot :

 melanoma

>
```

You may also receive a warning:

```
  The following object(s) are masked _by_ .GlobalEnv :

  read.geoda
```

which is because the `read.geoda` function you just created conflicts with the function of the same name that is included in the latest version of `spdep` .[1] If you did not source `read.geoda.R`, you will not have the warning. Note that the functions to read spatial weights have undergone several changes since the earlier *Data* tutorial was written. Specifically, the revised `read.gal` function can now read new-style GAL weights files, that is, files created by *GeoDa*. This workbook uses the most current version of these functions.

2.4.1 Using read.gal

Check out the features and parameters of the new `read.gal` function with the `help(read.gal)` command. You will see that there are three parameters: the name of the input file, a `region.id` and an `override.id`. The main purpose of these parameters is to avoid having to assume that the data set (the data frame) and the weights file have their observations in the same order. In *GeoDa*, the `ID` variable creates a key that can be matched with a corresponding variable in a data set to ensure the correct order. This is not an issue when both the data set and the weights are created together by *GeoDa*, as you just did, but it can be a problem when this is not the case.

For our purposes, it will suffice to set the `override.id = TRUE`, rather than the default `FALSE` to force `spdep` to use the new *GeoDa* GAL weights format. For example, with the `policerook.GAL` file as input, this becomes:[2]

```
> polgal <- read.gal("policerook.GAL",override.id=TRUE)
```

---

[1]These functions are basically the same, except that the "official" version has `skip = 0` as the default. So, to use the `read.geoda` function from `spdep`, make sure to set the option `skip = 1`.

[2]Make sure to check the spelling of the file name, specifically, whether the GAL file extension is in caps or lowercase. `R` is case sensitive.

This creates `polgal` as a neighbor list object. Technically, this is a list of lists, one for each observation, containing a list of the neighbors. In fact, it consists of more than just the neighbor information, and contains a list of ID values, and some metadata, as an `attributes` command reveals:

```
> attributes(polgal)
$class
[1] "nb"

$region.id
 [1] "28003" "28141" "28139" "28009" "28033" "28093" "28143" "28137" "28117"
[10] "28145" "28071" "28107" "28027" "28081" "28119" "28057" "28115" "28161"
[19] "28135" "28013" "28011" "28095" "28017" "28133" "28043" "28083" "28025"
[28] "28087" "28155" "28015" "28097" "28105" "28151" "28019" "28051" "28053"
[37] "28007" "28103" "28159" "28125" "28163" "28055" "28099" "28079" "28069"
[46] "28089" "28123" "28149" "28121" "28075" "28101" "28049" "28061" "28023"
[55] "28129" "28021" "28127" "28029" "28153" "28063" "28067" "28031" "28065"
[64] "28077" "28001" "28085" "28037" "28091" "28073" "28035" "28111" "28041"
[73] "28157" "28005" "28113" "28147" "28109" "28039" "28131" "28059" "28047"
[82] "28045"

$GeoDa
$GeoDa$shpfile
[1] "police"

$GeoDa$ind
[1] "FIPSNO"


$gal
[1] TRUE

$call
[1] TRUE

$sym
[1] TRUE
```

When you `print` the neighbor list, you no longer get a complete listing of its contents, as was the case in earlier versions. Instead, a summary-like overview of its main characteristics is listed:

```
> print(polgal)
Neighbour list object:
Number of regions: 82
Number of nonzero links: 402
Percentage nonzero weights: 5.978584
Average number of links: 4.902439
```

This may change again in the future.

2.4.2 Weights Characteristics

The characteristics of the weights file are obtained with the usual `summary` command:

```
> summary(polgal)
Neighbour list object:
Number of regions: 82
Number of nonzero links: 402
Percentage nonzero weights: 5.978584
Average number of links: 4.902439
Link number distribution:

 2  3  4  5  6  7
 1 13 16 23 21  8
1 least connected region:
28045 with 2 links
8 most connected regions:
28145 28071 28135 28043 28155 28007 28163 28085 with 7 links
```

Note how the link number distribution and summary of the connectedness structure use the `region.id` values to identify the observations. If you use `Tools > Weights > Properties` in *GeoDa*, you will observe the same frequency distribution in the link histogram.

2.5 Creating a Neighbor List from a GWT File

The new function `read.gwt2nb` included in the latest versions of `spdep` (included since the *Data* tutorial) is designed to convert a GWT weights file into a neighbor list object.

2.5.1 Using read.gwt2nb

The function `read.gwt2nb` takes as arguments the file name for the GWT file and the name of an `ID` variable, `region.id`. The latter must be "available" in your workspace. This is achieved by "attaching" the data frame that contains the variable. In our case, this is the `police` data frame. By attaching the data frame, the variables can be accessed by their name, e.g., as `FIPSNO`, instead of the more cumbersome `police$FIPSNO`. Attach the `police` data frame as:

```
> attach(police)
```

Make sure it worked by trying a `summary` of a variable, for example:

```
> summary(POLICE)
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
   49.0   291.5   432.5   927.8   955.8 10970.0
```

Now, turn the `policek5.GWT` file into a neighbor list, using `FIPSNO` as the `region.id`, and followed by a `summary`, as:

```
> polgwt <- read.gwt2nb("policek5.GWT",region.id=FIPSNO)
> summary(polgwt)
Neighbour list object:
Number of regions: 82
Number of nonzero links: 410
Percentage nonzero weights: 6.097561
Average number of links: 5
Non-symmetric neighbours list
Link number distribution:

 5
82
82 least connected regions:
28003 28141 28139 28009 28033 28093 28143 28137
... 28039 28131 28059 28047 28045 with 5 links
82 most connected regions:
28003 28141 28139 28009 28033 28093 28143 28137
... 28039 28131 28059 28047 28045 with 5 links
```

2.5.2 Checking for Symmetry of a Neighbor List

A useful function is a check for the symmetry of the neighbor list. Spatial weights based on contiguity should be symmetric, whereas weights derived from a k-nearest neighbor relation are typically not symmetric.

The `is.symmetric.nb` function in `spdep` implements this check. It takes the name of the neighbor object as argument. In our example, for the GAL-based weights:

```
> print(is.symmetric.nb(polgal))
[1] TRUE
```

and for the GWT-based weights,

```
> print(is.symmetric.nb(polgwt))
Non matching contiguities: 2 and 14
Non matching contiguities: 3 and 4
Non matching contiguities: 6 and 17
...
Non matching contiguities: 82 and 79
Non-symmetric neighbours list
[1] FALSE
```

2.6   Practice

Try the creation of a data frame, GAL-based neighbor lists and GWT-based neighbor lists for one of the other polygon sample data sets, such as the St Louis homicide data (STL HOM). Alternatively, try it on your own data set.

Exercise 3

Spatial Autocorrelation Analysis in R

3.1   Objectives

The purpose of this exercise is to illustrate how descriptive spatial autocorrelation analysis can be carried out in the R package spdep. We will also begin exploring some basic programming techniques in R, and start using plotting commands to design customized graphs.

The examples below all use the built-in COLUMBUS data set from spdep. Before trying the commands, make sure spdep is loaded by typing:

```
> spdep()
[1] "spdep, version 0.4-2, 2007-03-26"
```

If spdep is not loaded, you will get an error message. Also, load the COLUMBUS data set and attach it so that you can refer to the variables by name:

```
> data(columbus)
> attach(columbus)
```

For the practice sessions, you will be using the data frame for the POLICE data set and the GAL based neighbor list constructed in Exercise 2.

3.2   Converting Neighbor Lists to Spatial Weights

The neighbor list (an object of class nb) considered in Exercise 2 is only one of several classes that handle spatial contiguity information in spdep. It is

primarily intended to store information on the `ID` of the neighbor for each observation. For the spatial autocorrelation tests and spatial diagnostics in regression analysis, a different type of object is used, referred to as a *spatial weights object* (an object of class `listw`).

One converts an `nb` object into a `listw` object by means of the `nb2listw` function. This function has several options, and the defaults are not necessarily what one wants. An excerpt of `help(nb2listw)` shows the parameter list and the default settings:

```
Usage:

    nb2listw(neighbours, glist=NULL, style="W", zero.policy=FALSE)

Arguments:

neighbours: an object of class 'nb'
glist: list of general weights corresponding to neighbours
style: 'style' can take values W, B, C, U, and S
zero.policy: If FALSE stop with error for any empty neighbour sets, if
          TRUE permit the weights list to be formed with zero-length
          weights vectors
```

Two things to watch out for is to make sure that the "islands" are handled appropriately, since this affects future analysis. The default is to have any islands induce missing value codes through `zero.policy=FALSE`, which will stop the program when islands are encountered. To allow islands, make sure to explicitly set `zero.policy=TRUE`. As before, if you are happy with the default settings, there is no need to specify any of the parameters, and a simple `nb2listw(neighbor object)` will do.

3.2.1 Example

In the built-in COLUMBUS data set, the object `col.gal.nb` is a neighbor list for queen contiguity.[1] Since there are no islands in Columbus and the default of row-standardization (`style="W"`) is fine, we only need to specify the neighbor list name:

```
> colqueen <- nb2listw(col.gal.nb)
```

---

[1] Note that the queen contiguity weights are the ones originally included with the COLUMBUS data set on the now deprecated `geog55.geog.uiuc.edu` ftp site. These weights differ slightly from the ones constructed by *GeoDa*.

A check on the `class` of `colqueen` reveals that it is indeed an object of `class listw`:

```
> class(colqueen)
[1] "listw" "nb"
```

and `summary` gives the same connectedness information as before for the `nb` object. However, note how several summary measures of the weights have been precomputed, such as the various matrix traces used in the test statistics, more precisely, `nn, S0, S1` and `S2`:

```
> summary(colqueen)
Characteristics of weights list object:
Neighbour list object:
Number of regions: 49
Number of nonzero links: 230
Percentage nonzero weights: 9.579342
Average number of links: 4.693878
Link number distribution:

 2  3  4  5  6  7  8  9 10
 7  7 13  4  9  6  1  1  1
7 least connected regions:
1005 1008 1045 1047 1049 1048 1015 with 2 links
1 most connected region:
1017 with 10 links

Weights style: W
Weights constants summary:
   n   nn S0       S1        S2
W 49 2401 49 23.48489 204.6687
```

Finally, to see the actual weights contained in the object, use `colqueen$weights`, which is a list of lists of spatial weights for each observation:

```
> colqueen$weights
[[1]]
[1] 0.5 0.5

[[2]]
[1] 0.3333333 0.3333333 0.3333333
```

```
[[3]]
[1] 0.25 0.25 0.25 0.25
...
```

3.2.2 Practice

Use the `polgal` neighbor list object created in Exercise 2 for the POLICE
data set to turn it into a `listw` object and check its characteristics. If you
did not save `polgal` in your work space, you will need to recreate it from
the `policerook.GAL` file (follow the directions in in Section 2.4.1 on p. 8).
The results are listed below.

Creating a spatial weights list

```
> polrook <- nb2listw(polgal)
> summary(polrook)
Characteristics of weights list object:
Neighbour list object:
Number of regions: 82
Number of nonzero links: 402
Percentage nonzero weights: 5.978584
Average number of links: 4.902439
Link number distribution:

 2  3  4  5  6  7
 1 13 16 23 21  8
1 least connected region:
28045 with 2 links
8 most connected regions:
28145 28071 28135 28043 28155 28007 28163 28085 with 7 links

Weights style: W
Weights constants summary:
   n   nn S0       S1        S2
W 82 6724 82 35.14864 332.3231
> polrook$weights
[[1]]
[1] 0.3333333 0.3333333 0.3333333

[[2]]
```

```
[1] 0.3333333 0.3333333 0.3333333

[[3]]
[1] 0.25 0.25 0.25 0.25
...
```

3.3   Moran's I

Moran's I test statistic for spatial autocorrelation is implemented in `spdep`. There are two separate functions, `moran.test`, where inference is based on a normal or randomization assumption, and `moran.mc`, for a permutation-based test. Both take a variable name or numeric vector and a spatial weights list object (`listw`), in that order, as mandatory parameters. The permutation test also requires the number of permutations as a third mandatory parameter. Both functions also have several options, discussed in more detail below.

3.3.1   Normal and Randomization Inference

The different parameters and options for `moran.test` are revealed by a call to `help(moran.test)`:

```
moran.test(x, listw, randomisation=TRUE, zero.policy=FALSE,
           alternative="greater", rank = FALSE,
           na.action=na.fail, spChk=NULL)


Arguments:

x: a numeric vector the same length as the neighbours list in
    listw
listw: a 'listw' object created for example by 'nb2listw'
randomisation: variance of I calculated under the assumption of
          randomisation, if FALSE normality
zero.policy: if TRUE assign zero to the lagged value of zones
          without neighbours, if FALSE assign NA
alternative: a character string specifying the alternative
          hypothesis, must be one of greater (default), less
          or two.sided.
rank: logical value - default FALSE for continuous variables,
          if TRUE, uses the adaptation of Moran's I for ranks
```

```
          suggested by Cliff and Ord (1981, p. 46)
na.action: a function (default na.fail), can also be na.omit or
           na.exclude - in these cases the weights list will
           be subsetted to remove NAs in the data. It may be
           necessary to set zero.policy to TRUE
           because this subsetting may create no-neighbour
           observations.
           Note that only weights lists created without using
           the glist argument to nb2listw may be subsetted.
           If na.pass is used, zero is substituted for NA
           values in calculating the spatial lag
spChk: should the data vector names be checked against the
       spatial objects for identity integrity, TRUE, or FALSE,
       default NULL to use 'get.spChkOption()'
```

Of the optional parameters, two are very important. The `randomisation` option (watch the spelling!) is set to `TRUE` by default, which implies that in order to get inference based on a normal approximation, it must be explicitly set to `FALSE`. Similarly, the default is a *one-sided* test, so that in order to obtain the results for the (more commonly used) two-sided test, the option `alternative` must be set explicitly to `"two.sided"`. Note also that the `zero.policy` option is set to `FALSE` by default, which means that islands result in a missing value code (`NA`). Setting this option to `TRUE` will set the spatial lag for islands to the customary zero value.

To illustrate this, use the variable `CRIME` and the weights list `colqueen` with the normal approximation in a two-sided test:[2]

```
> moran.test(CRIME,colqueen,randomisation=FALSE,
+  alternative="two.sided")


        Moran's I test under normality

data:  CRIME
weights: colqueen

Moran I statistic standard deviate = 5.3818, p-value = 7.374e-08
alternative hypothesis: two.sided
sample estimates:
```

---

[2]Make sure you have the `columbus` data set attached, so that the variables can be referred to by name.

```
Moran I statistic        Expectation           Variance
     0.485770914       -0.020833333        0.008860962
```

Note how, unlike previous practice, the "object" created by `moran.test` was not assigned to a specific variable. If you simply want to get the results, this is not necessary. By entering the test this way, you indirectly invoke the `print` function for the object. If you want to access the individual results of the test for further processing, you should assign the `moran.test` to an object and then print that object. For example, using INC:

```
> mornorINC <- moran.test(INC,colqueen,randomisation=FALSE,
+ alternative="two.sided")
> print(mornorINC)


        Moran's I test under normality

data:  INC
weights: colqueen

Moran I statistic standard deviate = 4.6495, p-value = 3.327e-06
alternative hypothesis: two.sided
sample estimates:
Moran I statistic        Expectation           Variance
     0.416837942       -0.020833333        0.008860962
```

Note how the expectation and variance of the Moran's I test statistic is the same for both `CRIME` and `INC`. In fact, under the normal approximation, these moments only depend on the spatial weights, and not on the variable under consideration.

The `mornorINC` object belongs to the class `htest`, which is a generic class in R designed to hold the results of test statistics. Technically, it is a list and all its items can be accessed individually (provided you know what they are called; check the help). For example,

```
> class(mornorINC)
[1] "htest"
> mornorINC$statistic
Moran I statistic standard deviate
                          4.649514
```

shows the information stored in the `statistic` element of the list.

When the null hypothesis considered is based on the randomization distribution, the `randomisation` option does not need to be set. For example, again using CRIME:

```
> moran.test(CRIME,colqueen,alternative="two.sided")

        Moran's I test under randomisation

data:  CRIME
weights: colqueen

Moran I statistic standard deviate = 5.3427, p-value = 9.157e-08
alternative hypothesis: two.sided
sample estimates:
Moran I statistic       Expectation          Variance
     0.485770914      -0.020833333        0.008991121
```

Note how the value of the statistic and its expectation do not change relative to the normal case, only the variance is different (and thus the z-value and associated p-value).

Finally, leaving all the defaults as they are:

```
> moran.test(CRIME,colqueen)

        Moran's I test under randomisation

data:  CRIME
weights: colqueen

Moran I statistic standard deviate = 5.3427, p-value = 4.578e-08
alternative hypothesis: greater
sample estimates:
Moran I statistic       Expectation          Variance
     0.485770914      -0.020833333        0.008991121
```

where the only thing different is the p-value (half the p-value of the two-sided test).

3.3.2 Permutation Inference

A Moran's I test statistic with inference based on random permutation is contained in the function `moran.mc`. As in `moran.test`, it takes the variable

name (or vector) and the weights list file (`listw` object) as the first two mandatory arguments. It also needs the number of permutations as the third argument (`nsim`). Since the rank of the observed statistic is computed relative to the reference distribution of statistics for the permuted data sets, it is good practice to set this number to something ending on 9 (such as 99 or 999). This will yield nicely rounded pseudo p-values (like 0.01 or 0.001). The other options are identical to those for the `moran.test` function.

For example, using the CRIME variable with the `colqueen` weights, setting the number of permutations to 99 and leaving everything else to the default (e.g., a one-sided test):

```
> morpermCRIME <- moran.mc(CRIME,colqueen,99)
> morpermCRIME


        Monte-Carlo simulation of Moran's I

data:  CRIME
weights: colqueen
number of simulations + 1: 100


statistic = 0.4858, observed rank = 100, p-value = 0.01
alternative hypothesis: greater
```

Note how the number of simulations + 1 is a round 100. None of the permuted data sets yielded a Moran's I greater than the observed value of 0.4858, hence a pseudo p-value of 0.01.[3]

In this example, the results of the procedure were assigned to an object of class `htest` (`morpermCRIME`), which must be printed to reveal its contents. One item in the list `morpermCRIME` is a vector with the computed statistics for each of the permuted data set. This is contained in the item `$res`, as in:

```
> morpermCRIME$res
[1] -0.1211582450  0.0250296894 ... -0.0192339594 -0.1944962763
[6]  0.0397742260 -0.0554032219 ...  0.1625140071 -0.1785841236
   ...
[91]  0.0196008271 -0.1764712321 ... -0.0832054684  0.0782020854
[96]  0.0910332773 -0.1760530540 ... -0.0806383188  0.4857709137
```

Note that the "observed" value of Moran's I (0.48577) is included in this list as the last element. By default, the contents of this vector will slightly differ

---

[3]This is obtained as the ratio of the number of values greater or equal to observed statistic + 1 over the number of simulations + 1. In the example, this is $(0+1)/(99+1)=0.01$.

from run to run, since they are based on random permutations. The default random number seed value is determined from the current time and so no random permuation will be identical. To control the seed, use the R function `set.seed(seed, kind = NULL)` right before invoking the `moran.mc` command, and set the same value each time. For example, try this using `set.seed(123456)`:

```
> set.seed(123456)
> morpermCRIME <- moran.mc(CRIME,colqueen,99)
> morpermCRIME$res
  [1] -6.607790e-02  1.398150e-02  7.363572e-02  ...
  ...
  [100]  4.857709e-01
```

Unless you want to replicate results exactly, it is typically better to let the randomness of the clock set the seed.

3.3.3 Plotting the Reference Distribution

The full vector of Moran statistics for the permuted data sets lends itself well to a histogram or density plot. R has very sophisticated plotting functions, and this example will only scratch the surface. Consult the help files for further details and specific options.

In order to construct a density plot, we first need to create a "density" object. To make sure that this is the reference distribution plotted for the randomly permuted data sets only, we must remove the last element from `morpermCRIME$res`:

```
> morp <- morpermCRIME$res[1:length(morpermCRIME$res)-1]
```

Next, we must pass the new list (`morp`) to the function `density`. This function has many options, and in this example they are all kept to their default settings (check out `help(density)` for further details), so that only the vector with statistics is specified. The result is stored in the object `zz`:

```
> zz <- density(morp)
```

Next, we will plot three graphs on top of each other: a (continuous) density function (based on `zz`), a histogram for the reference distribution, and a line indicating the observed Moran's I. The latter is contained in the `statistic` attribute of the `moran.mc` object, `morpermCRIME$statistic`.

22

In the code that follows, the density plot is drawn first, along with the titles (`main` on top and `xlab` under the x-axis).[4] Also, to make the plot more distinct, a double line width is set (`lwd=2`) and the plot is in red (`col=2`). Finally, to make sure that the default settings for the plot do not make it too small, we explicitly set the horizontal (`xlim`) and vertical (`ylim`) width (typically, this takes some trial and error to get it right):

```
> plot.density(zz,main="Moran's I Permutation Test",
+ xlab="Reference Distribution",xlim=c(-0.3,0.7),
+ ylim=c(0,6),lwd=2,col=2)
```

Next, you add the histogram (`hist`) and a vertical line (`abline`). The histogram is drawn in regular thickness black lines, while the vertical line is double thickness (`lwd=2`) and drawn in blue (`col=4`):

```
> hist(morp,freq=F,add=T)
> abline(v=morpermCRIME$statistic,lwd=2,col=4)
```

The result is as in Figure 3.1 on p. 24.

Note that the example as described will plot the figure on your screen. In order to create hard copy output, you need to specify an output `device`. R supports several devices, but by far the most commonly used are the `postscript` and `pdf` devices. You must first open the device explicitly. For example,

```
>postscript(file="filename")
```

for a postscript file, or

```
>pdf(file="filename")
```

for a pdf file. All plotting commands that follow will be written to that device, until you close it with a `dev.off()` command. For example,

```
> postscript(file="moran.pdf")    # file name for pdf file
...                               # plotting commands
> dev.off()                       # close postscript device
```

---

[4]Several, but not all the plotting functions in R support an `add=T` argument, which adds a graph to an existing plot. Since the current version of `plot.density` does not support this argument, while the other two plots do, it has to be drawn first.

Figure 3.1: Moran's I Permutation Test Plot.

3.3.4 Practice

Use the POLICE data set and the `polrook` weights to test for spatial auto-correlation in the CRIME variable for the Mississippi counties. Use both the normal and randomization assumptions, as well as a permutation approach. Try making a graph of the distribution of the statistic in the simulated data sets, as in Figure 3.1.

Make sure to `detach(columbus)` and `attach(police)` before carrying out the tests so that you use the correct variable. Afterwards, do the reverse to get back to the COLUMBUS data for the remainder of the illustrations.

The results for the statistics are given below.

Moran's I, Normal Assumption

```
> moran.test(CRIME,polrook,randomisation=FALSE,
+ alternative="two.sided")

        Moran's I test under normality

data:  CRIME
weights: polrook

Moran I statistic standard deviate = 1.9108, p-value = 0.05604
alternative hypothesis: two.sided
sample estimates:
Moran I statistic        Expectation            Variance
     0.121668149        -0.012345679         0.004919118
```

Moran's I, Randomization Assumption

```
> moran.test(CRIME,polrook,alternative="two.sided")

        Moran's I test under randomisation

data:  CRIME
weights: polrook

Moran I statistic standard deviate = 2.4522, p-value = 0.01420
alternative hypothesis: two.sided
sample estimates:
Moran I statistic        Expectation            Variance
     0.121668149        -0.012345679         0.002986579
```

Moran's I, Permutation Approach

```
> morMissCRIME <- moran.mc(CRIME,polrook,999)
> print(morMissCRIME)

        Monte-Carlo simulation of Moran's I

data:  CRIME
```

```
weights: polrook
number of simulations + 1: 1000

statistic = 0.1217, observed rank = 986, p-value = 0.014
alternative hypothesis: greater
```

3.4   Constructing a Spatially Lagged Variable

Spatially lagged variables are important elements of many spatial regression specifications. In `spdep`, they are constructed by means of the `lag.listw` function. This function takes as arguments a spatial weights object of class `listw` and a conforming matrix (i.e., a matrix with the same number of rows as the dimension of the spatial weights).[5] Note that the documentation for `lag.listw` only refers to a "numeric vector," but this is incomplete. A look at the source code (type `lag.listw` without the usual parentheses) reveals that it works equally well for a matrix, i.e., for several variables in one function call. The third (optional) argument to `lag.listw` is the `zero.policy`, which we can ignore for the COLUMBUS and POLICE sample data sets.

3.4.1   Example

First, we will combine the three variables `CRIME`, `INC` and `HOVAL` into a matrix.[6] This is accomplished by means of the `cbind` command and will be an often used procedure when you program your own routines based on matrix algebra. For example, you create the matrix `xx` as:

```
> xx <- cbind(CRIME,INC,HOVAL)
```

and check the contents with

```
> xx
           CRIME     INC  HOVAL
 [1,] 15.725980 19.531 80.467
 [2,] 18.801754 21.232 44.567
 [3,] 30.626781 15.956 26.350
 ...
```

---

[5]The order of these arguments is the *reverse* of that in the Moran test statistics. Here, the weights object must come first, and the vector or matrix second. This order *must* be respected.

[6]Make sure you detached the `police` data frame and attached `columbus`.

Now, you can construct a spatially lagged version of this matrix, using the `colqueen` weights:

```
> wxx <- lag.listw(colqueen,xx)
```

and check the contents

```
> wxx
          [,1]      [,2]      [,3]
 [1,] 24.71427 18.594000 35.45850
 [2,] 26.24684 13.321333 46.67233
 [3,] 29.41175 14.123000 45.36475
 ...
```

Since this is a fairly small data set, we can easily check if the spatial lag terms are what they should be. Extract the neighbor IDs and spatial weights from `listw` by referring to the proper list elements (a `listw` object is a list, i.e., a combination of different elements). You see the names of the elements as the result of an `attributes` command (this lists the "attributes" of an object). For example,

```
> attributes(colqueen)
$names
[1] "style"      "neighbours" "weights"

$class
[1] "listw" "nb"

$region.id
 [1] 1005 1001 1006 1002 ... 1037 1039 1040 1009
...
[46] 1048 1015 1027 1026

$call
nb2listw(neighbours = col.gal.nb)
```

This shows that the neighbor IDs are in `colqueen$neighbours` (spelling!) and the corresponding weights in `colqueen$weights`. Specifically, for the first observation:

```
> colqueen$neighbours[1]
[[1]]
[1] 2 3
```

and

```
> colqueen$weights[1]
[[1]]
[1] 0.5 0.5
```

A simple calculation yields the spatial lags for the first observation "by hand":

```
> (xx[2,] + xx[3,])/2.0
   CRIME     INC    HOVAL
24.71427 18.59400 35.45850
```

which are the same values as in the first row of wxx.

Compare the descriptive characteristics of the original variables to their spatial lags:

```
> summary(xx)
    CRIME              INC              HOVAL
 Min.   : 0.1783   Min.   : 4.477   Min.   :17.90
 1st Qu.:20.0485   1st Qu.: 9.963   1st Qu.:25.70
 Median :34.0008   Median :13.380   Median :33.50
 Mean   :35.1288   Mean   :14.375   Mean   :38.44
 3rd Qu.:48.5855   3rd Qu.:18.324   3rd Qu.:43.30
 Max.   :68.8920   Max.   :31.070   Max.   :96.40
> summary(wxx)
       X1               X2               X3
 Min.   :13.85    Min.   : 8.539   Min.   :23.41
 1st Qu.:24.71    1st Qu.:10.957   1st Qu.:30.82
 Median :35.90    Median :14.303   Median :36.38
 Mean   :34.88    Mean   :14.742   Mean   :38.73
 3rd Qu.:45.39    3rd Qu.:17.161   3rd Qu.:45.36
 Max.   :54.91    Max.   :26.066   Max.   :75.13
```

Note how the spatial lag operation compresses the distribution of the variable. The measures of "spread" are substantially smaller than for the originals.

3.4.2 Practice

Compute the spatial lags for the CRIME, POLICE and INC variables in the police data set, using the polrook spatial weights. Compare the descriptive statistics of the original variables to their spatial lags. Experiment

with other spatial weights and/or try some statistical plots to visualize the difference between the distributions.

The basic results are given below.

Spatial Lags

```
> policex <- cbind(CRIME,POLICE,INC)
> wpolicex <- lag.listw(polrook,policex)
> summary(policex)
     CRIME            POLICE              INC
 Min.   :    5.0   Min.   :    49.0   Min.   : 4836
 1st Qu.:  76.5   1st Qu.:   291.5   1st Qu.: 6520
 Median :  141.5   Median :   432.5   Median : 7008
 Mean   :  194.5   Mean   :   927.8   Mean   : 7198
 3rd Qu.:  298.0   3rd Qu.:   955.8   3rd Qu.: 7612
 Max.   : 1739.0   Max.   : 10971.0   Max.   :10506
> summary(wpolicex)
       X1               X2                X3
 Min.   : 38.5   Min.   :  225.0   Min.   :6109
 1st Qu.:125.4   1st Qu.:  469.0   1st Qu.:6847
 Median :163.0   Median :  690.5   Median :7092
 Mean   :191.8   Mean   :  931.5   Mean   :7191
 3rd Qu.:224.6   3rd Qu.:  991.1   3rd Qu.:7449
 Max.   :732.7   Max.   : 3434.0   Max.   :8608
```

3.5   Moran Scatter Plot

The Moran scatter plot is also included in the `spdep` functionality. It is implemented in the function `moran.plot`. This takes as arguments a variable name, *followed by* a spatial weights objects of class `listw`. Optional parameters are the `zero.policy` as well as several ways to fine tune the labeling of high influence points, the x and y-axes, and the various graphical parameters (passed as in `par(..)`). The function `moran.plot` also creates a so-called "influence" object from the `influence.measures` passed back from the regression model that is used to estimate the slope (Moran's I). Setting the option `quiet=TRUE` suppresses the influence measures (the default is `quiet=FALSE`).

Figure 3.2: Moran Scatter Plot for `CRIME`, using `colqueen`.

3.5.1 Example

Use the `CRIME` variable from the `columbus` data frame and the spatial weights `colqueen` to invoke the Moran scatterplot, leaving all other settings to their default:

```
> moran.plot(CRIME,colqueen)
Potentially influential observations of
        lm(formula = wx ~ x) :

    dfb.1_   dfb.x   dffit   cov.r   cook.d   hat
```

```
1004  1.55_*  -1.40_*  1.55_*  0.59_*  0.86_*  0.11
1034  0.15    -0.21    -0.23   1.14_*  0.03    0.11
```

Note the list of influence characteristics and the IDs of the influential observations. The plot is as in Figure 3.2 on p. 30.

The Moran scatterplot produced by spdep differs somewhat from the one implemented in *GeoDa*. Most importantly, spdep uses the original variables, whereas *GeoDa* constructs the plot for standardized variates. Also, it is not clear how Moran's I can be extracted from the plot without an additional computation. The object returned by the function is a matrix with the observations as rows and the six influence measures as columns.

3.5.2 Customizing the Moran Scatter Plot

We will program a simple function that constructs a Moran scatter plot using standardized variables and shows the value of Moran's I on the plot. We will use the spdep function `lag.listw` to build the spatial lags and the generic R regression function `lm` to get the coefficients. First we test the various commands by trying them step by step, with hard-coded variables, then we generalize and turn them into a function.

First set the generic variable `x` to the variable of interest (CRIME), and convert it to a standardized variate, with a simple check to make sure we did it right:

```
> x <- CRIME
> zx <- (x - mean(x))/sd(x)
> mean(zx)
[1] -1.948555e-16
> var(zx)
[1] 1
```

Next, create the spatial lag for the standardized variable, after turning the weights into a generic variable:

```
> wfile <- colqueen
> wzx <- lag.listw(wfile,zx)
```

Now, compute the intercept and slope (Moran's I), using `lm`:

```
> morlm <- lm(wzx ~ zx)
> aa <- morlm$coefficients[1]
> mori <- morlm$coefficients[2]
```

Figure 3.3: Moran Scatter Plot for standardized `CRIME`, using `colqueen`.

```
> aa
(Intercept)
-0.01496451
> mori
        zx
0.4857709
```

Finally, we set up the plot as a simple scatterplot (`plot(x,y)`), add the regression line and the two axes (using `abline`), label the axes (using `xlab` and `ylab`), and put a title containing the value of Moran's I at the top

(using `title`, the `paste` command to put together text fragments, and a combination of `format` and `round` to make sure the Moran's I doesn't have too many decimal digits):

```
> par(pty="s")     # make sure it's square
> plot(zx,wzx,xlab="CRIME",ylab="Spatial Lag of CRIME")
> abline(aa,mori)
> abline(h=0,lty=2)
> abline(v=0,lty=2)
> title(paste("Moran Scatterplot I= ",format(round(mori,4))))
```

The result looks like Figure 3.3 on p. 32.

To combine these commands into a function, create a file in any text editor and name it `moran.plot2.R`. Enter the commands listed below and save the file as a text file in your working directory.

```
# moran.plot2
# function for Moran scatterplot using standardized variates
# specify variable = x
# specify spatial weights (listw object) = wfile
# all other options are hard coded (so far)
# example: moran.plot2(CRIME,colqueen)
#

moran.plot2 <- function(x,wfile)
{
    xname <- deparse(substitute(x))    # get name of variable
    zx <- (x - mean(x))/sd(x)
    wzx <- lag.listw(wfile,zx)
    morlm <- lm(wzx ~ zx)
    aa <- morlm$coefficients[1]
    mori <- morlm$coefficients[2]
    par(pty="s")
    plot(zx,wzx,xlab=xname,ylab=paste("Spatial Lag of ",xname))
    abline(aa,mori,col=2)
    abline(h=0,lty=2,col=4)
    abline(v=0,lty=2,col=4)
    title(paste("Moran Scatterplot I= ",format(round(mori,4))))
}
```

The only things different from before are the use of the functions `deparse`

and `substitute` to extract the variable name for use in the plot labels `xlab` and `ylab`. Source the function as

```
> source("moran.plot2.R")
```

and invoke it as `moran.plot2`. For example, passing `CRIME` and `colqueen` as the parameters:

```
> moran.plot2(CRIME,colqueen)
```

will yield the same plot as in Figure 3.3 (except that the regression line is in red and the axes drawn in blue, see Figure 3.4 on p. 35).

3.5.3 Practice

Try out the new function using the `police` data frame (make sure to first `detach(columbus)` then `attach(police)`), and construct a Moran scatterplot for the variable `INC` using `polrook` as the weight file. Compare this plot to the standard `spdep` `moran.plot`. Try some other variables and/or spatial weights. Also try to get your plot written out to a postscript or pdf file (make sure to include `dev.off()` at the end to switch back to screen output).

If you feel adventurous, check out the source code of `moran.plot` and add the various safety checks and the diagnostics of high leverage to `moran.plot2`.

Moran Scatter Plot

```
> pdf(file="morscatf.pdf")
> moran.plot2(INC,polrook)
> dev.off()
```

Moran Scatterplot I= 0.0991

Figure 3.4: Moran Scatter Plot for standardized `INC`, using `polrook`.

Exercise 4

Monte Carlo Simulation (1) Assessing the Properties of a Test Statistic

4.1 Objectives

The purpose of this exercise is to illustrate how to set up a simple Monte Carlo simulation experiment to assess some properties of a test statistic. Specifically, the end goal is to describe the distribution and compare the nominal and empirical rejection frequencies for Moran's I under the null hypothesis. In the process, we will look at how to generate random variables and carry out basic matrix algebra in R. We will also use **spdep** to create spatial weights for regular grids and convert spatial weights to actual matrix objects.

As before, make sure **spdep** is loaded before starting the exercises. Since all the data will be generated, there is no need to specify a particular data frame.

4.2 Generating Random Variables

R has extensive functionality to generate random variables for a wide range of distributions and to compute probabilities for those distributions.[1]

Each distribution function has an R name, such as **norm** for the normal and **unif** for the uniform. The different random variable related functions

---

[1] See Chapter 8 on Probability Distributions in Venables et al. (2004) for further details.

are indicated by a single letter prefix to the distribution name. This includes `d` for the density, `p` for the cumulative density function, `q` for the quantile function, and `r` for simulation of random variables.[2] The arguments to these functions are, respectively, the value of a random variable `x` for `dxxx` and `pxxx`, a probability for `qxxx`, and the sample size for `rxxx`.

4.2.1 Example

To generate a vector of 25 standard normal variates, use `rnorm` with only the sample size as parameter, as in:

```
> x <- rnorm(25)
```

Check the mean and standard deviation:[3]

```
> mean(x)
[1] -0.199015
> sd(x)
[1] 0.9869352
```

Note how for a sample this small, the mean is fairly far from its theoretical expectation. To get a much closer fit and eliminate some of the extra variability induced by the random number generator, you need to generate much larger sets. For example:

```
> x <- rnorm(100000)
> mean(x)
[1] -0.005083976
> sd(x)
[1] 0.996822
> x <- rnorm(1000000)
> mean(x)
[1] -0.001698689
> sd(x)
[1] 1.000263
```

To generate a vector of normal random variates with a given mean and standard deviation, you need to pass these as arguments in addition to the

---

[2]The CDF and quantile functions have the special arguments `lower.tail` and `log` that are very useful in practice. For example, to compute the upper tail of a distribution as when the p-value of a statistic is required, or, 1 - CDF(x), set `lower.tail = FALSE`.

[3]Note that your results will vary, since the seed of the random number generator was not fixed using `set.seed`.

sample size. Note that you pass the *standard deviation*, **not** the variance. For example:

```
> x <- rnorm(1000,5,20)
> mean(x)
[1] 5.241835
> sd(x)
[1] 19.59694
```

You can now compare a given percentile in the generated distribution to its theoretical value. For example, in the 1000 generated values above the 95th percentile will be in row 950 of the sorted vector `x`:

```
> z <- sort(x)
> z[950]
[1] 36.46587
```

The theoretical value is found with the `qnorm` function, passing 0.95 (the probability) as the argument, in addition to the mean and standard deviation:

```
> qz <- qnorm(0.95,5,20)
> qz
[1] 37.89707
```

A plot of the sorted z values against qz values for the same percentile would yield a q-q plot. Conversely, finding the cumulative probability for `z[950]`:

```
> pz <- pnorm(z[950],5,20)
> pz
[1] 0.9421746
```

4.2.2   Practice

Experiment with some of the other distributions, such as the uniform, using `runif(n,min,max)`, or the Poisson, using `rpois(n,lambda)`. Use the techniques from Exercise 3 to draw a histogram or density plot for the generated variates and summarize the characteristics of the distribution (`mean`, `sd`, and `summary`).

Figure 4.1: Histogram for Randomly Generated Poisson Variates.

Histogram for Poisson Variates

```
> z2 <- rpois(50,2)
> summary(z2)
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
   0.00    1.00    2.00    2.14    3.00    6.00
> mean(z2)
[1] 2.14
> var(z2)
[1] 2.041224
```

```
> pdf(file="poisson.pdf")
> hist(z2)
> dev.off()
```

4.3   Basic Matrix Algebra Operations

Operations on matrices and vectors are element by element in R. So, if `x` is a vector (say `n` by 1) and `A` is a `n` by `n` matrix, then `A * x` will multiply each row element in `A` with the matching row element in `x` and yield another `n` by n matrix, not a n by 1 vector. To illustrate, use the `c` operator (for concatenate) to create a 5 by 1 vector `x` (by default this is a column vector), and the `matrix` function with row and column parameters to create a 5 by 5 matrix (of sequence numbers):

```
> x <- c(1:5)
> x
[1] 1 2 3 4 5
> A <- matrix(1:25,5,5)
> A
     [,1] [,2] [,3] [,4] [,5]
[1,]    1    6   11   16   21
[2,]    2    7   12   17   22
[3,]    3    8   13   18   23
[4,]    4    9   14   19   24
[5,]    5   10   15   20   25
```

Note how the default ordering of a matrix is in so-called "column major order" (down each column first). Alternatively, a vector can be passed to the `array` function with a `dim` attribute for the same effect (this is useful for reshaping matrices as well):

```
> z <- c(1:25)
> B <- array(z,dim=c(5,5))
> B
     [,1] [,2] [,3] [,4] [,5]
[1,]    1    6   11   16   21
[2,]    2    7   12   17   22
[3,]    3    8   13   18   23
[4,]    4    9   14   19   24
[5,]    5   10   15   20   25
```

Now, multiply `A` by `x`:

```
> zz <- A * x
> zz
     [,1] [,2] [,3] [,4] [,5]
[1,]    1    6   11   16   21
[2,]    4   14   24   34   44
[3,]    9   24   39   54   69
[4,]   16   36   56   76   96
[5,]   25   50   75  100  125
```

In order to get *true* matrix multiplication, following to the rules of matrix algebra, a different operator must be used, symbolized as `%*%`:

```
> zz1 <- A %*% x
> zz1
     [,1]
[1,]  215
[2,]  230
[3,]  245
[4,]  260
[5,]  275
```

A few other important basic matrix operations are:

- `nrow(matrix)` for the number of rows

- `ncol(matrix)` for the number of columns

- `t(matrix)` to construct the transpose of a matrix, e.g., X'

- `diag(matrix)` to extract the diagonal elements from a matrix

- `diag(vector)` to create a diagonal matrix with `vector` as diagonal

- `crossprod(matrix1,matrix2)` to compute the crossproduct (e.g., $X'y$) between two matrices efficiently

- `%o%` for the outer product of two vectors, e.g., xx'

- `outer(vector1,vector2,function)` for any function applied to each pair x_i and y_j

41

4.3.1 Example

Start with a vector of 30 standard normal variates and convert it to a 10 by 3 matrix. Note how you can force "row major order" by setting the argument `byrow=TRUE`:

```
> x1 <- rnorm(30)
> x1
 [1]  1.86263657 -0.14248903 -1.26807899  ... -0.30481949
 [7]  1.05654040  0.46209724 -0.75313876  ... -1.95010825
[13] -0.33947593  0.55823272 -1.42072261  ...  0.37998187
[19] -0.75149464  0.47130272  0.76668664  ...  0.64171921
[25] -0.99857864 -1.53179763  1.94797343  ... -0.32623853
> XX <- matrix(x1,10,3)
> XX
             [,1]        [,2]        [,3]
 [1,]   1.8626366  1.02988875  0.7666866
 [2,]  -0.1424890 -1.95010825  0.3450489
 [3,]  -1.2680790 -0.33947593  0.6046427
 [4,]   0.2091366  0.55823272  0.6417192
 [5,]   0.4177341 -1.42072261 -0.9985786
 [6,]  -0.3048195  0.06214715 -1.5317976
 [7,]   1.0565404 -0.31582345  1.9479734
 [8,]   0.4620972  0.37998187 -0.9578326
 [9,]  -0.7531388 -0.75149464 -1.5308449
[10,]   0.7728531  0.47130272 -0.3262385
> XX <- matrix(x1,10,3,byrow=TRUE)
> XX
             [,1]        [,2]        [,3]
 [1,]   1.86263657 -0.1424890 -1.2680790
 [2,]   0.20913657  0.4177341 -0.3048195
 [3,]   1.05654040  0.4620972 -0.7531388
 [4,]   0.77285309  1.0298887 -1.9501082
 [5,]  -0.33947593  0.5582327 -1.4207226
 [6,]   0.06214715 -0.3158235  0.3799819
 [7,]  -0.75149464  0.4713027  0.7666866
 [8,]   0.34504894  0.6046427  0.6417192
 [9,]  -0.99857864 -1.5317976  1.9479734
[10,]  -0.95783257 -1.5308449 -0.3262385
```

Check the number of rows and columns and construct the transpose of matrix `XX`:

```
> rXX <- nrow(XX)
> rXX
[1] 10
> cXX <- ncol(XX)
> cXX
[1] 3
> tXX <- t(XX)
> tXX
           [,1]       [,2]       [,3]  ...          [,6]
[1,]  1.8626366  0.2091366  1.0565404  ...   0.06214715
[2,] -0.1424890  0.4177341  0.4620972  ...  -0.31582345
[3,] -1.2680790 -0.3048195 -0.7531388  ...   0.37998187
           [,7]      [,8]       [,9]      [,10]
[1,] -0.7514946 0.3450489 -0.9985786 -0.9578326
[2,]  0.4713027 0.6046427 -1.5317976 -1.5308449
[3,]  0.7666866 0.6417192  1.9479734 -0.3262385
```

Compute the cross product of the two matrices (e.g., $X'X$), extract the diagonal and divide by the number of rows. If these were in deviations from the mean, the result would be the estimated variance in each column of the original matrix. Note how the division / is an element by element operation:

```
> XtX <- crossprod(XX)
> XtX
           [,1]       [,2]       [,3]
[1,]  7.944244   3.747365 -6.210131
[2,]  3.747365   7.157987 -4.951293
[3,] -6.210131 -4.951293 13.134562
> v1 <- diag(XtX)
> v1
[1]  7.944244  7.157987 13.134562
> var1 <- v1 / nrow(XX)
> var1
[1] 0.7944244 0.7157987 1.3134562
```

4.3.2 Practice

Make sure you become familiar with the difference between element by element operations and true matrix algebra in R. Practice by creating random

43

vectors and matrices, multiply them, compute cross-products and extract diagonals. These will be important tools in constructing the test statistics needed in the simulation experiment.

4.4 Creating Spatial Weights for a Grid Layout

In many simulation settings, one controls for the spatial layout of the data by generating a sample for a regular square grid. The weights matrix for such a layout has a simple structure and does not require the checking of boundaries or computation of distances that is needed for irregular layouts. The function `cell2nb` in `spdep` creates a neighbor object for any rectangular grid, using either rook or queen as the criterion. Alternatively, a so-called torus correction may be used to avoid edge effects.[4]

The `cell2nb` function takes as required parameters the number of rows and the number of columns in the grid. In addition, the type of contiguity can be specified (default is `type="rook"`, the other option is `queen`), as well as whether or not a torus correction is needed (set `torus=TRUE`, the default is `FALSE`).

4.4.1 Example

To create a rook type neighbor object for a 4 by 4 square grid layout, and to check its properties, use:

```
> rook4x4 <- cell2nb(4,4)
> summary(rook4x4)
Neighbour list object:
Number of regions: 16
Number of nonzero links: 48
Percentage nonzero weights: 18.75
Average number of links: 3
Link number distribution:

2 3 4
4 8 4
4 least connected regions:
1:1 4:1 1:4 4:4 with 2 links
```

---

[4]The torus correction connects the grid cells in the right-most cell to those in the left-most cell, and those in the top row to those in the bottom row, thus ensuring that each cell has four neighbors (for rook).

```
4 most connected regions:
2:2 3:2 2:3 3:3 with 4 links
```

Compare to the properties with a torus correction:

```
> rook4x4t <- cell2nb(4,4,torus=TRUE)
> summary(rook4x4t)
Neighbour list object:
Number of regions: 16
Number of nonzero links: 64
Percentage nonzero weights: 25
Average number of links: 4
Link number distribution:

 4
16
16 least connected regions:
1:1 2:1 3:1 4:1 1:2 2:2 3:2 4:2 ... 1:4 2:4 3:4 4:4 with 4 links
16 most connected regions:
1:1 2:1 3:1 4:1 1:2 2:2 3:2 4:2 ... 1:4 2:4 3:4 4:4 with 4 links
```

4.4.2 Practice

Create neighbor objects for a 5 by 10 rectangular grid using both rook and queen as contiguity criterion, and compare their properties. Experiment with other layouts and/or a torus correction.

Creating Rectangular Grid Neighbor Objects

```
> rook5x10 <- cell2nb(5,10)
> summary(rook5x10)
Neighbour list object:
Number of regions: 50
Number of nonzero links: 170
Percentage nonzero weights: 6.8
Average number of links: 3.4
Link number distribution:

 2  3  4
 4 22 24
4 least connected regions:
```

45

```
1:1 5:1 1:10 5:10 with 2 links
24 most connected regions:
2:2 3:2 4:2 ... 4:9 with 4 links
> queen5x10 <- cell2nb(5,10,type="queen")
> summary(queen5x10)
Neighbour list object:
Number of regions: 50
Number of nonzero links: 314
Percentage nonzero weights: 12.56
Average number of links: 6.28
Link number distribution:

 3  5  8
 4 22 24
4 least connected regions:
1:1 5:1 1:10 5:10 with 3 links
24 most connected regions:
2:2 3:2 4:2 ... 4:9 with 8 links
```

4.5   Converting Spatial Weights to a Matrix

The neighbor list `nb` objects and the spatial weights `listw` objects both store the contiguity information in an efficient (sparse) form. Sometime, it is necessary to use the full matrix representation of the weights, for example, to use them in matrix computations.[5]

The function `nb2mat` converts the contiguity information in an `nb` object to a full matrix. A call to the `help` function shows the necessary arguments and default options:

```
> help(nb2mat)
nb2mat                     package:spdep                     R Documentation

Spatial weights matrices for neighbours lists

Description:

     The function generates a weights matrix for a neighbours list with
```

---

[5]Typically, however, it is *very inefficient* to operate on the full matrix, and sparse operations should be pursued as much as possible.

```
          spatial weights for the chosen coding scheme.

Usage:

     nb2mat(neighbours, glist=NULL, style="W", zero.policy=FALSE)
     listw2mat(listw)

Arguments:

neighbours: an object of class 'nb'

   glist: list of general weights corresponding to neighbours

   style: 'style' can take values W, B, C, and S

zero.policy: If FALSE stop with error for any empty neighbour sets, if
          TRUE permit the weights list to be formed with zero-length
          weights vectors

   listw: a 'listw' object from for example 'nb2listw'
```

4.5.1 Example

Convert the previously created rook neighbor list for the 4 by 4 regular grid
to a full matrix, using row-standardized weights, as:

```
> wrks <- nb2mat(rook4x4)
> wrks
            [,1]      [,2]      [,3]      [,4]      [,5]      [,6]      [,7]
 [1,] 0.0000000 0.5000000 0.0000000 0.0000000 0.5000000 0.0000000 0.0000000
 [2,] 0.3333333 0.0000000 0.3333333 0.0000000 0.0000000 0.3333333 0.0000000
 [3,] 0.0000000 0.3333333 0.0000000 0.3333333 0.0000000 0.0000000 0.3333333
 [4,] 0.0000000 0.0000000 0.5000000 0.0000000 0.0000000 0.0000000 0.0000000
 [5,] 0.3333333 0.0000000 0.0000000 0.0000000 0.0000000 0.3333333 0.0000000
 [6,] 0.0000000 0.2500000 0.0000000 0.0000000 0.2500000 0.0000000 0.2500000
 [7,] 0.0000000 0.0000000 0.2500000 0.0000000 0.0000000 0.2500000 0.0000000
 [8,] 0.0000000 0.0000000 0.0000000 0.3333333 0.0000000 0.0000000 0.3333333
 [9,] 0.0000000 0.0000000 0.0000000 0.0000000 0.3333333 0.0000000 0.0000000
[10,] 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.2500000 0.0000000
[11,] 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.2500000
```

```
[12,] 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
[13,] 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
[14,] 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
[15,] 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
[16,] 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
            [,8]      [,9]     [,10]     [,11]     [,12]     [,13]     [,14]
 [1,] 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
 [2,] 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
 [3,] 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
 [4,] 0.5000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
 [5,] 0.0000000 0.3333333 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
 [6,] 0.0000000 0.0000000 0.2500000 0.0000000 0.0000000 0.0000000 0.0000000
 [7,] 0.2500000 0.0000000 0.0000000 0.2500000 0.0000000 0.0000000 0.0000000
 [8,] 0.0000000 0.0000000 0.0000000 0.0000000 0.3333333 0.0000000 0.0000000
 [9,] 0.0000000 0.0000000 0.3333333 0.0000000 0.0000000 0.3333333 0.0000000
[10,] 0.0000000 0.2500000 0.0000000 0.2500000 0.0000000 0.0000000 0.2500000
[11,] 0.0000000 0.0000000 0.2500000 0.0000000 0.2500000 0.0000000 0.0000000
[12,] 0.3333333 0.0000000 0.0000000 0.3333333 0.0000000 0.0000000 0.0000000
[13,] 0.0000000 0.5000000 0.0000000 0.0000000 0.0000000 0.0000000 0.5000000
[14,] 0.0000000 0.0000000 0.3333333 0.0000000 0.0000000 0.3333333 0.0000000
[15,] 0.0000000 0.0000000 0.0000000 0.3333333 0.0000000 0.0000000 0.3333333
[16,] 0.0000000 0.0000000 0.0000000 0.0000000 0.5000000 0.0000000 0.0000000
           [,15]     [,16]
 [1,] 0.0000000 0.0000000
 [2,] 0.0000000 0.0000000
 [3,] 0.0000000 0.0000000
 [4,] 0.0000000 0.0000000
 [5,] 0.0000000 0.0000000
 [6,] 0.0000000 0.0000000
 [7,] 0.0000000 0.0000000
 [8,] 0.0000000 0.0000000
 [9,] 0.0000000 0.0000000
[10,] 0.0000000 0.0000000
[11,] 0.2500000 0.0000000
[12,] 0.0000000 0.3333333
[13,] 0.0000000 0.0000000
[14,] 0.3333333 0.0000000
[15,] 0.0000000 0.3333333
[16,] 0.5000000 0.0000000
attr(,"call")
```

```
nb2mat(neighbours = rook4x4)
```

4.5.2   Practice

Create a matrix from the queen or other neighbor objects you constructed in 4.4.2. Use a matrix multiplication operation to construct a spatially lagged variable for a conforming random vector. Compare the result to what the `lag.listw` function yields (you must make sure to convert the `nb` object to a `listw` using `nb2listw`).

Creating a Spatial Lag Using Matrix Multiplication

```
> x <- rnorm(50)
> w50 <- nb2mat(queen5x10)
> w50[1,]
 [1] 0.0000000 0.3333333 0.0000000 0.0000000 0.0000000 0.3333333 0.3333333 ...
> wx <- w50 %*% x
> wwx <- lag.listw(nb2listw(queen5x10),x)
> wx[1:4]
        1:1         2:1         3:1         4:1
 0.0765760 -0.2601768   0.1325916 -0.5860503
> wwx[1:4]
[1]  0.0765760 -0.2601768   0.1325916 -0.5860503
```

4.6   A Simulation Experiment

An important property of any test statistic is its distribution under the null, and, more specifically, characteristics such as the mean, variance, and higher moments. In addition, the *size* of the test statistic is the probability of rejecting the null for a given significance level. For a test statistic to be *unbiased*, its size should equal the p-value corresponding to the theoretical critical value (the nominal significance). For example, if 1.96 is chose as the critical value in a two-sided test for a statistic that follows a standard normal distribution, 5% of the cases should reject the null hypothesis, even though it is true.

A simple simulation experiment can demonstrate these properties for Moran's I, applied to a generic random variable.[6] We will start with considering uncorrelated standard normal variates, but this can easily be extended to other distributions.

_____

[6]The theoretical moments for Moran's I applied to regression residuals are **different**.

4.6.1   Computing Moran's I Using Matrix Algebra

For an $n \times n$ row-standardized weights matrix W with elements w_{ij}, the Moran's I statistic is:

$$I = z'Wz/z'z, \tag{4.1}$$

where z is a $n \times 1$ vector of variables expressed as deviations from the mean. When the weights are not row-standardized, a scaling factor of n/S_0 is applied, where $S_0 = \sum_i \sum_j w_{ij}$.

In our simulation experiment, we will not be relying on the built-in Moran's I calculations from spdep, but instead compute the statistic directly using matrix algebra and the spatial lag operation. The first step in this process is to generate a random vector and convert it to deviations from the mean:

```
> y <- rnorm(16)
> yz <- y - mean(y)
> mean(yz)
[1] -4.163336e-17
```

Next, we build the numerator and denominator of equation (4.1), using matrix products and the spatial lag operator `lag.listw` with the rook weights for the grid:

```
> yy <- crossprod(yz)
> yy
          [,1]
[1,] 13.98298
> wy <- lag.listw(nb2listw(rook4x4),yz)
> ywy <- crossprod(yz,wy)
> ywy
          [,1]
[1,] 1.294643
```

The Moran's I statistic follows directly, as:

```
> mi <- ywy / yy
> mi
            [,1]
[1,] 0.09258704
```

4.6.2 The Theoretical Moments of Moran's I

Under the normal assumption for the null, the theoretical moments of Moran's I only depend on the characteristics of the weights matrix. Specifically,

$$E[I] = \frac{-1}{n-1} \tag{4.2}$$

$$E[I^2] = \frac{n^2 S_1 - n S_2 + 3 S_0^2}{(n-1)(n+1) S_0^2}, \tag{4.3}$$

where S_1 and S_2 are sums of weights matrix elements,

$$S_1 = (1/2) \sum_i \sum_j (w_{ij} + w_{ji})^2 \tag{4.4}$$

$$S_2 = \sum_i (w_{i*} + w_{*j})^2, \tag{4.5}$$

with w_{i*} as the row sums and w_{*j} as the column sums. Note that with row-standardized weights, these expressions simplify, since each row sum equals 1, and therefore $S_0 = n$.

The variance of Moran's I then follows as

$$Var[I] = E[I^2] - E[I]^2, \tag{4.6}$$

and the statistic is turned into a standardized variate as

$$I_z = \frac{I - E[I]}{\sqrt{Var[I]}}. \tag{4.7}$$

Since the mean and variance of the statistic under the null depend solely on the weights matrix, and not on the actual variables under consideration, they must only be computed once. The required "sums" can be extracted from a spatial weights object by means of the **spdep** function `spweights.constants`.

Consider the `rook4x4` neighbor object for the square grid and extract the sums by converting the `nb` object to a `listw` object (using `nb2listw`):

```
> rks <- spweights.constants(nb2listw(rook4x4))
> rks
$n
[1] 16
$n1
[1] 15
```

```
$n2
[1] 14
$n3
[1] 13
$nn
[1] 256
$S0
[1] 16
$S1
[1] 11.05556
$S2
[1] 64.61111
```

Using these results, we can compute the moments of Moran's I in this example:

```
> eirk4x4 <- -1 / rks$n1
> eirk4x4
[1] -0.06666667
> ei2rk4x4 <- (rks$n^2 * rks$S1 - rks$n * rks$S2 + 3 * rks$S0^2)/
+ (rks$n1 * (rks$n + 1) * rks$S0^2)
> ei2rk4x4
[1] 0.03928377
> varirk4x4 <- ei2rk4x4 - eirk4x4^2
> varirk4x4
[1] 0.03483932
> sdirk4x4 <- sqrt(varirk4x4)
> sdirk4x4
[1] 0.1866530
```

4.6.3   Inference

Inference is based on a normal approximation, where the statistic is standardized and then compared to the standard normal distribution.

We standardize the statistic computed in section 4.6.1 with the moments computed above:

```
> iz <- (mi - eirk4x4)/sdirk4x4
> iz
          [,1]
[1,] 0.8532076
```

and obtain a two-tailed significance using `pnorm` with the `lower.tail = FALSE` option, and multiplying the result by 2.0

```
> piz <- pnorm(iz,lower.tail=FALSE) * 2
> piz
          [,1]
[1,] 0.3935442
```

To check this, we can use the random vector `y` in the `moran.test` with the rook grid weights (see section 3.3.1 on p. 17):

```
> moran.test(y,nb2listw(rook4x4),randomisation=FALSE,
+ alternative="two.sided")


        Moran's I test under normality

data:  y
weights: nb2listw(rook4x4)

Moran I statistic standard deviate = 0.8532, p-value = 0.3935
alternative hypothesis: two.sided
sample estimates:
Moran I statistic       Expectation          Variance
      0.09258704       -0.06666667        0.03483932
```

These results are identical to those computed "by hand."

4.6.4 Setting up the Experiment

The simulation experiment will require the following steps:

- initialize the number of simulations (`r`), sample size (`n`) and critical value (either a p-value, such as 0.05, or its counterpart in the distribution, such as 1.96)

- create a neighbor object for a given lattice layout (see `cell2nb` in section 4.4)

- convert the neighbor object to a weights list (`nb2listw`)

- compute the theoretical moments of Moran's I (section 4.6.2)

- initialize vectors with results (value of statistic, decision)

53

- loop over each random replication

  - generate standard normal vector (section 4.2)
  - compute Moran's I (section 4.6.1) and store result in vector
  - compute p-value (section 4.6.3) and store decision in vector

- summarize results

  - compute mean and variance (standard deviation) of results vector with Moran's I
  - do tests on normal distribution for results vector (q-q plot)
  - plot the density of the statistics, compute proportion above 1.96
  - compute percent rejection

4.6.5 The Simulation Loop

The simulation loop requires so-called "flow control." We will use a simple `for` loop, which iterates over elements in a sequence. For example, we will compute 1000 iterations of our random vectors and associated statistics.[7] We start by initializing the number of replications as well as the two results vectors:

```
> r <- 1000
> moran <- vector(mode="numeric",length=r)
> reject <- vector(mode="numeric",length=r)
```

Also, to avoid having to convert the `rook4x4` neighbor list to a `listw` object over and over again, we create a new spatial weights object, and also set the sample size to 16 and the critical value to `crit`:

```
> n <- 16
> crit <- 0.05
> w4 <- nb2listw(rook4x4)
```

Next, we organize the main computations into a loop, using `for(i in 1:r)` as the iterator:

---

[7]Using 1000 replications is for illustrative purposes only, a "real" simulation experiment would require a much larger sample to achieve a reasonable degree of precision.

```
> for (i in 1:r){
+     y <- rnorm(n)
+     yz <- y - mean(y)
+     yy <- crossprod(yz)
+     wy <- lag.listw(w4,yz)
+     ywy <- crossprod(yz,wy)
+     mi <- ywy / yy
+     morans[i] <- mi
+     iz <- (mi - eirk4x4)/sdirk4x4
+     piz <- pnorm(iz,lower.tail=FALSE) * 2
+     if (piz < crit) { reject[i] <- 1 }
+ }
```

Note the use of an `if` statement in the loop. The condition in parentheses is checked, and if it is evaluated as `TRUE` the action that follows is executed. In this example, this means that when the null hypothesis is rejected (as a two-sided test with a p-value of 0.05) then it is stored in the `reject` vector. Otherwise, the initialized value of zero remains.

4.6.6 Analyzing the Results

The results are contained in the vectors `morans` and `reject`. We can use the contents of `morans` to get the summary descriptive statistics for the Moran's I under the simulated null hypothesis, and use it to create histograms and density plots.

   For example, the mean and variance in the simulation were (these results will obviously vary from simulation to simulation):

```
> mean(morans)
[1] -0.0744135
> sd(morans)
[1] 0.185987
> summary(morans)
    Min.  1st Qu.   Median     Mean  3rd Qu.     Max.
-0.75700 -0.19800 -0.07318 -0.07441  0.04712  0.50160
```

These estimated moments can now be compared to the theoretical results obtained in section 4.6.2.

   The rejection frequency is the sum of rejections (values of 1 in the vector `reject`) divided by the number of replications (`r`):

```
> rejfreq <- sum(reject) / r
> rejfreq
[1] 0.018
```

While this seems well below the nominal Type I error of `crit = 0.05`, a proper comparison needs to take into account the random variability in the simulation itself. The rejection frequency can be considered to be an estimate for the underlying probability of a series of Bernoulli random variates, each with $p = 0.05$. The variance associated with r such draws is $p(1-p)/r$. With $p = 0.05$ and $r = 1000$, this variance and the corresponding standard error of the simulation experiment are:

```
> varsim <- crit * (1.0 - crit) / r
> varsim
[1] 4.75e-05
> sdsim <- sqrt(varsim)
> sdsim
[1] 0.006892024
```

An interval around the true Type I error of two standard deviations would then be:

```
> simint <- c(crit - 2*sdsim,crit + 2*sdsim)
> simint
[1] 0.03621595 0.06378405
```

Taking this into account, our experiment would suggest that the Moran's I test statistic "significantly" under-rejects the null hypothesis in a small data set of 16 observations, using the rook criterion to define contiguity.

4.6.7 Practice

You have now all the pieces to put together a function that simulates Moran's I under the null hypothesis for a generic sample size, number of replications and weights matrix, reporting at a minimum the mean, standard deviation and rejection frequency. Using the examples in Exercise 3, you can also create density plots or histograms for the simulated statistic. You can assess the difference in the properties of the test statistic between rook and queen definition of contiguity and compare to when a torus adjustment is made. You can experiment with distributions other than the normal as well.

As a challenge, you can also consider what it would take to compute the inference under the randomization assumption.

Simulating Moran's I

For example, consider the following small function to run the simulations for any rectangular lattice, using the rook contiguity and the randomization option to carry out inference. Note that this function is *slow*, since it uses the built-in `moran.test` function (and associated overhead) in each iteration. For simulations, considerable speed up can be obtained by limiting the computations purely to the necessary items, without the overhead associated with all the features of a function like `moran.test`.

```
# sim.moran.rand.R
# generates reference distribution for Moran's I
# calculates rejection frequency using randomization
# Usage:
#     sim.moran.rand(rn,cn,rpl)
# Arguments
#     rn: number of rows in the lattice
#     cn: number of columns in the lattice
#     rpl: number of replications

sim.moran.rand <- function(rn,cn,rpl,pvalue=0.05)
{
      morans <- vector(mode="numeric",length=rpl)
      reject <- vector(mode="numeric",length=rpl)
      w <- nb2listw(cell2nb(rn,cn))
      n <- rn * cn
      for (i in 1:rpl){
          y <- rnorm(n)
          mi <- moran.test(y,w,alternative="two.sided")
          morans[i] <- mi$statistic
          if (mi$p.value) { reject[i] <- 1}
          }
      rejfreq <- sum(reject) / rpi
      value <- list(morans = morans, rej = rejfreq)
 }
```

To run this function, source it first (`>source("sim.moran.rand.R")`), then invoke it, for example. as in:

```
> resu <- sim.moran.rand(5,5,1000)
```

The result is a list containing a vector with all the computed statistics in
`resu$morans` and the rejection frequency in `resu$rej`:

```
> resu$rej
[1] 0.06
```

A pdf file with a frequency plot of the simulated distribution is con-
structed as:

```
> pdf(file="moransim.pdf")
> plot.density(zz,main="Moran's I under Null - Randomization",
+ xlab="Reference Distribution",lwd=2,col=2)
> dev.off()
```

The resulting plot is given in Figure 4.2 on p. 59. Other summaries are
possible as well, such as histograms, qq plots or descriptive statistics.

Figure 4.2: Moran's I Under the Null.

Exercise 5

Monte Carlo Simulation (2) Assessing the Properties of an Estimator

5.1 Objectives

The purpose of this exercise is to develop some further skills in designing and carrying out a Monte Carlo simulation experiment, this time focused on the properties of an estimator. The main objective is to obtain results for the bias, standard error and mean squared error of an estimator for different scenarios of error distributions. Along the way, we will review ordinary least squares regression in R, as well as how to simulate spatially correlated error terms (using spdep).

As before, make sure spdep is loaded before starting the exercises, using `library(spdep)`, and that the `columbus` data frame is attached.

5.2 Ordinary Least Squares Regression

We will focus on the ordinary least squares estimator in the linear regression model:

$$y = X\beta + \epsilon, \tag{5.1}$$

where y is a n by 1 vector of observations on the dependent variable, X is a n by k matrix of observations on the explanatory variables with matching k by 1 coefficient vector β, and ϵ is a n by 1 vector of random error terms.

For now, we will assume that the error terms are independent, identically distributed (i.i.d.).

The familiar OLS estimator for β, $\hat{\beta}$ is:

$$\hat{\beta} = (X'X)^{-1}X'y, \qquad (5.2)$$

where $(X'X)$ is a k by k matrix and $X'y$ is a k by 1 vector.

Under the classical regularity conditions, the OLS estimator is unbiased, such that $E[\hat{\beta}] = \beta$ and its precision can be obtained from

$$Var[\hat{\beta}] = \sigma^2(X'X)^{-1}, \qquad (5.3)$$

where σ^2 is the *unknown* error variance, estimated as $\hat{\sigma^2} = e'e/(n-k)$, with $e = y - X\hat{\beta}$ as the vector of OLS residuals.

5.2.1 OLS Using lm

The standard approach to get OLS estimates in R is to use the *linear model* or `lm` functionality. This gets passed a `formula` object, which takes the form `y ~ x1 + x2`, where `y` is the dependent variable and the variables to the right of the $\sim$ are the explanatory variables. Check `help(lm)` for extensive documentation of the various options.

For example, using the `columbus` data frame and the variables `CRIME`, `INC`, and `HOVAL`, the standard regression yields:

```
> col1 <- lm(CRIME ~ INC + HOVAL)
> summary(col1)

Call:
lm(formula = CRIME ~ INC + HOVAL)

Residuals:
    Min      1Q  Median      3Q     Max
-34.418  -6.388  -1.580   9.052  28.649

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  68.6190     4.7355  14.490  < 2e-16 ***
INC          -1.5973     0.3341  -4.780 1.83e-05 ***
HOVAL        -0.2739     0.1032  -2.654   0.0109 *
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 11.43 on 46 degrees of freedom
Multiple R-Squared: 0.5524,    Adjusted R-squared: 0.5329
F-statistic: 28.39 on 2 and 46 DF,  p-value: 9.34e-09
```

The result object `col1` has many attributes, including the vector of coefficients, `col1$coefficients`, the vector of residuals, `col1$residuals`, and the predicted values, `col1$fitted.values`.

5.2.2   OLS Matrix Algebra

In a simulation experiment, one needs to avoid recomputing items that do not change between replications. For example, if an OLS estimation is part of a replication loop, and the design matrix (X) does not change between replications, using a call to `lm` is not a good idea. Instead, you want to compute all the parts that do not change *before* the loop is started and only keep what changes inside the loop.

In the case of OLS, this means that we need to compute the estimates by using matrix algebra. From equation 5.2, we know that this involves the computation of two cross products, the inverse of a cross product, and the product of a matrix with a vector. The matrix products are straightforward to implement, using the `crossprod` and `%*%` operators. The inverse is something that usually requires some care, and for numerical stability a Cholesky decomposition or qr approach is often preferable. For this example, we will use the `solve` operator. This supports two operations, the inverse of a matrix, such as `solve(A)`, and the *solution* to a system of linear equations, as `solve(A,b)`, where the result equals $A^{-1}b$.

5.2.3   Example

Before proceeding to the solution of the OLS normal equations, we first need to turn the relevant elements of the `columbus` data frame into vectors and matrices. For the dependent variable, this is a simple assignment:

```
> y <- CRIME
> y
 [1] 15.725980 18.801754 30.626781 ...  0.178269
 [8] 38.425858 30.515917 34.000835 ... 57.066132
...
[43] 36.663612 25.962263 29.028488 ... 22.541491
```

For the matrix of observations on the explanatory variables, X, we use the cbind function and exploit the recycling rule to add a vector of ones as the first column:

```
> X <- cbind(1,INC,HOVAL)
> X
          INC  HOVAL
 [1,] 1 19.531 80.467
 [2,] 1 21.232 44.567
...
[49,] 1 18.796 35.800
```

Note how even though we only include the *scalar* 1, the recycling rule turns this into a vector of the same length as the other two vectors.

Next, we need to construct the matrix of cross products $X'X$ and the vector $X'y$:

```
> XX <- crossprod(X)
> XX
                    INC      HOVAL
      49.000    704.372   1883.375
INC  704.372 11686.673 29600.442
HOVAL 1883.375 29600.442 88757.619
> Xy <- crossprod(X,y)
> Xy
         [,1]
      1721.312
INC   21557.532
HOVAL 57640.624
```

Note how we only need to specify X as the argument in the first crossprod operation, since R constructs $X'X$ by default. The inverse $(X'X)^{-1}$ follows directly, as:

```
> XXi <- solve(XX)
> XXi
                          INC          HOVAL
      0.171498009 -0.0072068055 -0.0012356172
INC   -0.007206805  0.0008538132 -0.0001318211
HOVAL -0.001235617 -0.0001318211  0.0000814476
> XX %*% XXi
```

```
                            INC         HOVAL
        1.000000e+00  3.018903e-17  7.543415e-18
  INC   -5.933489e-15  1.000000e+00 -1.250043e-16
  HOVAL -1.711690e-14 -1.422552e-15  1.000000e+00
```

The OLS estimates follow as:

```
> bols <- XXi %*% Xy
> bols
            [,1]
      68.6189611
INC   -1.5973108
HOVAL -0.2739315
```

A more efficient (and numerically more accurate) way is to include both `XX` and `Xy` in the `solve` operation, as:

```
> b <- solve(XX,Xy)
> b
            [,1]
      68.6189611
INC   -1.5973108
HOVAL -0.2739315
```

The drawback of this for our simulation is that it does not give us a separate $(X'X)^{-1}$, which is the part that remains constant between replications. Also, we need the inverse separately if we want to compute the standard errors of the OLS estimates and the usual t-test statistics.

First, we need to obtain an estimate for the error variance. This requires the calculation of the residuals and their sum of squares, which then need to be divided by the degrees of freedom, as:

```
> e <- y - X %*% b
> e

                [,1]
   [1,]    0.3465419
   [2,]   -3.6947990
...
  [49,]   -6.2476690
> ee <- crossprod(e)
> ee
```

```
         [,1]
[1,] 6014.893
> vare <- ee / (nrow(X) - ncol(X))
> vare

         [,1]
[1,] 130.7585
 > sqrt(vare)

         [,1]
[1,] 11.43497
```

Compare to the results printed in the summary of the `lm` function.

The standard errors of the estimates are the square roots of the diagonal elements of the matrix $\hat{\sigma}^2(X'X)^{-1}$. We extract these by means of the `diag` function and apply the square root function, to yield:

```
> varb <- vare * diag(XXi)
> seb <- sqrt(varb)
> seb
[1] 4.7354861 0.3341308 0.1031987
```

Note how we do not need to use the matrix multiplication operator to multiply the scalar `vare` with the diagonal elements of the inverse matrix. The t-test statistics then follow as the ratio of the coefficient estimates over their standard error:

```
> tb <- b / seb
> tb
           [,1]
      14.490373
INC   -4.780496
HOVAL -2.654409
```

The associated probabilities are found from the Student t distribution (note the use of the `abs` function for this to work properly):

```
> pb <- pt(abs(tb),df=nrow(X)-ncol(X),lower.tail=FALSE) * 2
> pb
             [,1]
      9.210890e-19
```

```
INC   1.828960e-05
HOVAL 1.087450e-02
```

Again, compare to the output of `lm` applied to the Columbus example.

5.2.4 Practice

Use the POLICE sample data set to replicate the regression in the 1992 *Regional Science and Urban Economics* Kelejian-Robinson article. For simplicity, you could also use a subset of the regressors. The Kelejian-Robinson specification was a regression of police expenditures on the tax rate, intergovernmental transfers, per capital income, crime rate, unemployment rate, home ownership, percent college, percent white and percent commuters (see Kelejian and Robinson 1992, pp. 323–324). Use both the standard `lm` function as well as an explicit matrix computation.

Kelejian-Robinson Baseline OLS Regression

```
> police.ols <- lm(POLICE ~ TAX + TRANSFER + INC + CRIME +
+ UNEMP + OWN + COLLEGE + WHITE + COMMUTE)
> summary(police.ols)

Call:
lm(formula = POLICE ~ TAX + TRANSFER + INC + CRIME + UNEMP +
    OWN + COLLEGE + WHITE + COMMUTE)

Residuals:
    Min     1Q  Median      3Q     Max
-841.75  -99.18   15.94  127.21 1007.11

Coefficients:
             Estimate Std. Error t value Pr(>|t|)
(Intercept) -4.858e+02  5.136e+02  -0.946   0.3474
TAX          2.182e-01  4.715e-01   0.463   0.6450
TRANSFER     7.553e-02  2.819e-03  26.789  < 2e-16 ***
INC          1.028e-01  5.031e-02   2.043   0.0447 *
CRIME        1.331e+00  1.661e-01   8.011 1.46e-11 ***
UNEMP       -1.990e+01  1.724e+01  -1.154   0.2522
OWN         -7.929e+00  5.658e+00  -1.401   0.1654
COLLEGE     -3.686e-01  4.013e+00  -0.092   0.9271
WHITE       -3.408e-01  2.586e+00  -0.132   0.8955
```

```
COMMUTE      3.268e+00  2.996e+00   1.091   0.2789
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 272.4 on 72 degrees of freedom
Multiple R-Squared: 0.9708, Adjusted R-squared: 0.9671
F-statistic: 265.9 on 9 and 72 DF,  p-value: < 2.2e-16
```

5.3   Spatial Autoregressive Random Variables

In Exercise 4, we conducted a simulation experiment under the null, when
it was not necessary to specify a particular form or magnitude for the spa-
tial autocorrelation. In order to obtain spatially autocorrelated random
variables, we will need to apply a transformation that turns uncorrelated
variates into correlated ones. We will focus on spatial autoregressive (SAR)
and spatial moving average (SMA) processes, although this by no means ex-
hausts the set of possibilities. In this section, we deal with SAR processes.
SMA processes are treated in section 5.4.

In its simplest form, a spatial autoregressive process is:

$$\epsilon = \rho W \epsilon + u \tag{5.4}$$

where u and ϵ are n by 1 random vectors, ρ is the autoregressive parameter,
and W is a n by n spatial weights matrix. The u vector is i.i.d., and ϵ is
spatially correlated. We can obtain ϵ directly from u by means of a so-called
spatial autoregressive transformation:

$$\epsilon = (I - \rho W)^{-1} u. \tag{5.5}$$

The package `spdep` contains the function `invIrM` which constructs the
inverse matrix $(I - \rho W)^{-1}$ required in a spatial autoregressive transforma-
tion. It takes as arguments a neighbor list (`nb`) or weights list (`listw`)
object, followed by a value for the autoregressive parameter (`rho`). In addi-
tion, a list of general weights can be optionally included, as well as a specific
style for the weights (the default is `style="W"`, a row-standardized weights
matrix). The result is an n by n matrix. In contrast to the original spatial
weights, which tend to be sparse, the inverse is typically dense, requiring a
lot of memory for larger data sets.

5.3.1 Example

Use the `col.gal.nb` neighbor list for Columbus to create a spatial autoregressive transformation matrix, with 0.5 for the autoregressive parameter:

```
> c05 <- invIrM(col.gal.nb,0.5)
```

Now use this transformation matrix to turn a vector of 49 standard (uncorrelated) random variates into a vector of spatially correlated random variables, as:[1]

```
> uu <- rnorm(49)
> mean(uu);sd(uu)
[1] 0.2138447
[1] 1.247263
> e05 <- c05 %*% uu
> mean(e05);sd(e05)
[1] 0.3113457
[1] 1.397368
```

Check with a Moran spatial autocorrelation test:

```
> moran.test(e05,nb2listw(col.gal.nb),randomisation=FALSE,
+ alternative="two.sided")

        Moran's I test under normality

data:  e05
weights: nb2listw(col.gal.nb)

Moran I statistic standard deviate = 3.4981, p-value = 0.0004687
alternative hypothesis: two.sided
sample estimates:
Moran I statistic       Expectation          Variance
     0.308448612      -0.020833333       0.008860962
```

5.3.2 Practice

Use one of the spatial weights created for the regular lattices in Exercise 4, or the spatial weights for the POLICE sample data set to generate vectors

---

[1]While the standard deviation of the spatially correlated `e05` is reported, this is not totally valid, since these variables are heteroskedastic and thus do not have a common variance.

with spatially autocorrelated normal random variables. Experiment with different values for the spatial autoregressive parameter and check if the `moran.test` rejects the null. Specifically, check for small values of `rho`, such as 0.1 or 0.05.

Creating Spatially Autoregressive Variables

Consider a simple function (without any error checking) that takes as arguments a vector, a `nb` object and a value for the autoregressive parameter. It uses this and the `spdep` function `invIrW` to create a spatial autoregressive transformation.

```
# sar.transform
# transform a vector into spatially autoregressive vector
# Usage:
#    sar.transform(x,w,rho)
# Arguments:
#    x: vector of values
#    w: spatial weights nb object
#    rho: spatial autoregressive coefficient
# Value:
#    vector
# uses spdep invIrM
# WARNING: no error checking!

sar.transform <- function(x,w,rho)
{
    irw <- invIrM(w,rho)
    wx <- irw %*% x
    wx
}
```

Create a spatial autoregressive transformation of the `POLICE` variable, and test for spatial autocorrelation as:

```
> source("sar.transform.R")
> wpolice05 <- sar.transform(POLICE,polgal,0.5)
> moran.test(wpolice05,polrook,alternative="two.sided")

Moran's I test under randomisation
```

```
data:  wpolice05
weights: polrook

Moran I statistic standard deviate = 6.1307, p-value = 8.752e-10
alternative hypothesis: two.sided
sample estimates:
Moran I statistic        Expectation          Variance
     0.365968180        -0.012345679        0.003807943
```

5.4   Spatial Moving Average Random Variables

A spatial moving average process is:

$$\epsilon = \rho W u + u. \tag{5.6}$$

In order to construct a random variable that follows this process, no inverse
is required, but only the sum of a random variable and its spatial lag, the
latter multiplied by the spatial parameter. This is a simple operation that
does not require matrix operations, but can be implemented using the sparse
`lag.listw` function in spdep.

5.4.1   Example

Using the same `col.gal.nb` neighbor list and the same vector of standard
normal variates `uu` as in Section 5.3.1, create a spatial lag for the random
vector:

```
> wu <- lag.listw(nb2listw(col.gal.nb),uu)
```

Now, carry out a spatial moving average transformation as:

```
> em05 <- 0.5 * wu + uu
> mean(em05);sd(em05)
[1] 0.2490343
[1] 1.279354
```

Compare the degree of spatial autocorrelation to that of the SAR coun-
terpart with a `moran.test`:

```
> moran.test(em05,nb2listw(col.gal.nb),randomisation=FALSE,
+ alternative="two.sided")
```

```
        Moran's I test under normality

data:  em05
weights: nb2listw(col.gal.nb)

Moran I statistic standard deviate = 2.5339, p-value = 0.01128
alternative hypothesis: two.sided
sample estimates:
Moran I statistic      Expectation         Variance
     0.217686051     -0.020833333      0.008860962
```

5.4.2   Practice

Use the same weights and normal random variates as in 5.3.2 to create a new
random vector that follows a spatial moving average process. Compare the
degree of spatial autocorrelation as indicated by a Moran test for different
values of the spatial parameter.

Creating Spatial Moving Average Variables

Again, consider a simple function (without any error checking) that takes as
arguments a vector, a **nb** object and a value for the spatial moving average
parameter. It uses this and the **spdep** function **lag.listw** to create a spatial
moving average transformation.

```
# sma.transform
# transform a vector into spatial moving average vector
# Usage:
#    sma.transform(x,w,rho)
# Arguments:
#    x: vector of values
#    w: spatial weights nb object
#    rho: spatial autoregressive coefficient
# Value:
#    vector
# uses spdep lag.listw
# WARNING: no error checking!

sma.transform <- function(x,w,rho)
{
    wx <- lag.listw(nb2listw(w),x)
```

```
    smax <- rho * wx + x
    smax
}
```

As before, source the function and invoke it to transform the `POLICE` variable into a moving average form.

```
> source("sma.transform.R")
> smapolice05 <- sma.transform(POLICE,polgal,0.5)
> moran.test(smapolice05,polrook,alternative="two.sided")

Moran's I test under randomisation

data:  smapolice05
weights: polrook

Moran I statistic standard deviate = 5.1891, p-value = 2.113e-07
alternative hypothesis: two.sided
sample estimates:
Moran I statistic        Expectation             Variance
      0.302526841       -0.012345679          0.003682022
```

5.5   A Simulation Experiment

We will implement a simple simulation experiment to assess the properties of the OLS estimator when the regression error term follows a spatial autoregressive process. We know from theory that OLS should remain unbiased, but will be less efficient than for i.i.d. errors. Also, the standard errors of the estimates as suggested by OLS will be biased and tend to underestimate the true variability of the estimates.

The basic design of the experiment consists of generating vectors of simulated "observations" on y by using $y = X\beta + \epsilon$, for a fixed matrix X and a given parameter vector β. The ϵ vectors are constructed by a spatial autoregressive transformation. Starting with an i.i.d. random vector u, a value for the parameter ρ and a spatial weights matrix, $\epsilon = (I - \rho W)^{-1}u$. The simulated ϵ is added to $X\beta$ to yield the y vector. The product $X\beta$ is fixed and must only be computed once, so it should be *outside* the main simulation loop.

There are many design aspects that need to be controlled for in a simulation experiment such as this one. Paradoxically, one aspect that requires

72

very little attention is the choice of the value for the *true* β parameters. Typically, one takes a single vector of random X values and sets $\beta = 1$ for simplicity. In order to make sure the regressions have a similar and reasonable fit, the variance of the simulated X is chosen together with the variance of the error term. Consider that the true regression total variance (of y) can be decomposed into variance due to X and variance due to the error term, we select the two variances such that the regression R^2 is something like 0.9. For example, we could randomly generate X as a normal variate with mean zero and variance 9, and the error term as a standard normal variate. This would yield regressions with "on average" an R^2 of 0.9. Note that this only holds under the null of no spatial autocorrelation. The spatially correlated error terms will no longer be homoskedastic, so that the notion of an R^2 is not valid in a strict sense, nor is there a common error variance.

The properties of the estimator can be quantified in a number of ways. First, we consider the *bias* of the estimator, which is estimated as the difference between the average of $\hat{\beta}$ across all replications and the true parameter value β:

$$bias = \sum_r \hat{\beta}_r / R - \beta \tag{5.7}$$

where R is the total number of replications. The precision of the estimator can be measured as the variance (or, standard error) of the reference distribution of the $\hat{\beta}_r$:

$$Var[\hat{\beta}] = (1/R) \sum_r (\hat{\beta}_r - \bar{\hat{\beta}})^2, \tag{5.8}$$

where $\bar{\hat{\beta}} = \sum_r \hat{\beta}_r / R$, and the standard error is the square root of this result. The result for the simulated data can be compared to its theoretical counterpart, the matching diagonal element in $\sigma^2 (X'X)^{-1}$ with σ^2 replaced by the value chosen for the simulated error term.

Alternatively, the measure of precision can be computed with respect to the *true* value of β, yielding the *mean squared error* (MSE):

$$MSE[\hat{\beta}] = (1/R) \sum_r (\hat{\beta}_r - \beta)^2. \tag{5.9}$$

The MSE is particularly appropriate when comparing estimators that are not necessarily unbiased, since it consists of the sum of the variance and the squared bias. For an unbiased estimator, the latter is zero ("in expected value") and MSE and variance are the same.

5.5.1 Setting up the Experiment

The simulation experiment will require the following steps:

- initialization of parameters

  - set the number of replications (`r`), sample size (`n`), value range for the autoregressive parameters (`rho`), variance of the error terms (`sigma`)

- initialization and preprocessing of X matrix

  - generate a random vector (`x`) with variance such that the regression R^2 is a given target value (e.g., 0.9).
  - create the `X` matrix as $1 + $ `x`, this implicitly sets β to (1.0, 1.0)
  - compute `XXi` $= (X'X)^{-1}$
  - compute `vb2` $= \sigma^2 (X'X)_{2,2}^{-1}$, the theoretical variance of OLS

- initialization and preprocessing of spatial weights

  - create a neighbor object for a chosen regular lattice layout (`cell2nb`)
  - if needed, convert neighbor object to list weights (`nb2listw`)

- initialize vector (or matrix) to hold estimated coefficients

- loop over values for `rho`

  - create SAR transformation matrix for given `rho` and weights (`invIrM`)

- main simulation loop

  - generate i.i.d. random vector `uu`
  - transform `uu` to spatially autoregressive vector `eps`
  - compute dependent variable vector `y` $=$ `Xb` $+$ `eps`
  - compute cross product `Xy` $= X'y$
  - obtain OLS estimate `b` $=$ `XXi %*% Xy`
  - store `b` in results vector

- summarize results

  - compute average `b`

74

- compute bias
- compute variance of `b`
- compare variance to "theoretical" variance
- compute MSE of `b`
- (optional) graph reference distribution of `b`

5.5.2   Example

We start with the various initializations, to keep things generic for later incorporation into a function:

```
> n <- 49
> r <- 1000
> rho <- c(0.0, 0.2, 0.5, 0.7, 0.9)
```

Next we construct the `X` matrix, its cross product and inverse:

```
> X <- cbind(1,rnorm(n,0,3))
> XX <- crossprod(X)
> XX
           [,1]        [,2]
[1,] 49.000000  -3.904727
[2,] -3.904727 391.827489
> XXi <- solve(XX)
> XXi
              [,1]           [,2]
[1,] 0.0204243828 0.0002035376
[2,] 0.0002035376 0.0025541719
> vb2 <- XXi[2,2]
> vb2
[1] 0.002554172
```

Note that we use 3 as the standard deviation in the standard normal variate (to obtain a variance of 9) and how again the recycling rule is exploited to create the constant term. The theoretical variance of the OLS estimator for β_2 is extracted as the `[2,2]` element of the variance matrix (since $\sigma^2 = 1$). Also, the value of $X\beta$ can be pre-computed as the row sum of the matrix `X` (rather than an explicit matrix-vector multiplication):

```
> Xb <- rowSums(X)
```

Now, we construct a neighbor list for rook contiguity on a 7 by 7 regular lattice. The spatial autoregressive transformation is different for each value of `rho`, so that must be carried out inside the loop over the values of `rho`:

```
> w <- cell2nb(7,7)
```

Next, initialize the matrix that will hold the results for β_2 (the slope) for each of the values of `rho` (0.0 to 0.9, as columns) and in each of the replications:

```
> bres <- matrix(0, nrow=r, ncol= length(rho))
> bres
        [,1] [,2] [,3] [,4] [,5]
   [1,]    0    0    0    0    0
...
 [999,]    0    0    0    0    0
[1000,]    0    0    0    0    0
```

We are now ready for the main simulation loop. There is actually a double loop, one over the elements of `rho`, the second over the replications. In the outer loop, the spatial autoregressive transformation matrix is built for the corresponding value of `rho`. In the inner loop, the random errors are generated and transformed, the "observed" `y` is constructed, the regression is estimated and the result for $\hat{\beta}$ stored in `bres`:

```
> for (j in 1:length(rho)) {
+        iw <- invIrM(w,rho[j])
+        for (i in 1:r) {
+            e <- iw %*% rnorm(n)
+            y <- Xb + e
+            Xy <- crossprod(X,y)
+            b <- XXi %*% Xy
+            bres[i,j] <- b[2]
+        }
+ }
```

The summary characteristics of the OLS estimator are:

```
> summary(bres)
      X1                 X2                 X3                 X4
 Min.   :0.8474    Min.   :0.8357    Min.   :0.8364    Min.   :0.8155
 1st Qu.:0.9670    1st Qu.:0.9663    1st Qu.:0.9640    1st Qu.:0.9637
```

```
Median :0.9996    Median :1.0015    Median :0.9986    Median :1.0013
Mean   :0.9993    Mean   :1.0006    Mean   :0.9999    Mean   :1.0006
3rd Qu.:1.0333    3rd Qu.:1.0361    3rd Qu.:1.0349    3rd Qu.:1.0392
Max.   :1.1515    Max.   :1.1707    Max.   :1.2232    Max.   :1.1838
     X5
Min.   :0.7475
1st Qu.:0.9337
Median :0.9990
Mean   :0.9985
3rd Qu.:1.0607
Max.   :1.2515
```

This illustrates how the mean and median don't change much from the expected value of 1, but the range of the distribution widens with ρ (first slightly, but then quite dramatically with $\rho = 0.9$). More precise characteristics of the distribution are the average (`bavg`, note the use of the function `colMeans`), the bias (`bias`), variance (`varb`, note again the use of `colMeans`), the standard error (`sdeb`), difference with the theoretical standard error (the square root of `vb2`), and the mean squared error (`mseb`):

```
> bavg <- colMeans(bres)
> bavg
[1] 0.9993119 1.0006317 0.9999488 1.0006297 0.9985044
> bias <- bavg - 1.0
> bias
[1] -6.880809e-04  6.316951e-04 -5.124518e-05  6.296576e-04
    -1.495599e-03
> varb <- colMeans((bres - bavg)^2)
> varb
[1] 0.002373828 0.002616419 0.002882452 0.003484847 0.008237334
> sdeb <- sqrt(varb)
> sdeb
[1] 0.04872195 0.05115094 0.05368848 0.05903260 0.09075976
> sdeb - sqrt(vb2)
[1] -0.0018168670  0.0006121246  0.0031496605  0.0084937801
     0.0402209418
> mseb <- colMeans((bres - 1)^2)
> mseb
[1] 0.002371185 0.002614124 0.002878571 0.003485413 0.008236709
```

In essence, the OLS estimator remains unbiased, even with $\rho = 0.9$, but its variance, standard error and mean squared error are affected. They increase

with ρ, first slightly, but then much more so for $\rho = 0.9$.

To end in beauty, we will plot the empirical densities for the estimates obtained in the simulation, for each of the values of ρ and use some fancy plotting techniques to combine them all in one graph, plot in different colors, and add a legend to boot. As in Exercise 3, we start by turning the columns of the matrix `bres` into `density` objects, `zz0` through `zz4`:

```
> zz0 <- density(bres[,1])
> zz1 <- density(bres[,2])
> zz2 <- density(bres[,3])
> zz3 <- density(bres[,4])
> zz4 <- density(bres[,5])
```

Next, we draw a density plot for the last graph (`zz4`), set the `main` title and the label for the x-axis (`xlab`), and specify the color to be red (2). In addition, to make sure that all the plots will fit on the same graph, we use the min and max values from the `summary(bres)` to set the `xlim` and `ylim` parameters (the range for the values on the x- and y-axis respectively). We further add the other graphs as `lines` (since the `add` option is not implemented for a `density.plot`), specifying the x-value as `zz $x` and the y-value as `zz $y`, with a different color for each. Finally, we add a `legend` (one of the more complex but also more powerful options). We set the coordinates for the upper left corner, specify the text for the legend (using `paste` to add the words `rho =` in front of the value taken from the `rho` vector), set the type to line (`lty=1`), in the same dimension as the legend "words" (i.e., a five element vector), and set the matching colors in a vector (`col`):

```
> plot.density(zz4,main="Density of beta hat",xlab="beta hat",
+ xlim=c(0.7,1.3),ylim=c(0,8),col=2)
> lines(zz3$x,zz3$y,col=6)
> lines(zz2$x,zz2$y,col=5)
> lines(zz1$x,zz1$y,col=3)
> lines(zz0$x,zz0$y,col=4)
> legend(0.7,8,paste("rho = ",rho),lty=c(1,1,1,1,1),
+ col=c(4,3,5,6,2))
```

The result looks like Figure 5.1, illustrating the effect of ρ on the spread of the distribution (of course, you need colors to distinguish the lines).

Density of beta hat

Figure 5.1: Empirical Distribution of $\hat{\beta}$ for different ρ.

5.5.3 Practice

You can now combine the various commands into a function that allows you to make the sample size generic, specify the type of distribution for the error term, include both spatial autoregressive as well as spatial moving average transformations, etc. In addition, you could try to assess the effect on OLS of an omitted spatial lag variable (the transformation is slightly more complex than in the error case). The sky is the limit ...

Exercise 6

Regression Diagnostics for Spatial Autocorrelation

6.1 Objectives

This exercise provides and overview of specification testing for spatial autocorrelation in regression residuals. We will build upon the basic linear regression reviewed in Section 5.2 and use as a running example the homicide regression for Southern US counties, detailed in Chapters 24 and 25 of the *GeoDa Workbook* (Anselin 2005). We will focus on Moran's I for regression residuals and on the Lagrange Multiplier test statistics for spatial error and spatial lag correlation.

6.2 Preliminaries

As before, make sure `spdep` is loaded before starting the exercises, using `library(spdep)`.

First, we will create two text files containing a subset of the variables from the SOUTH data set. For the purposes of this exercise, we will be using the variables pertaining to 1960 and 1990 and create a different data frame for each. Use the same approach as outlined in Section 2.2.1. Specifically, load the `south.shp` file into *GeoDa*, with `FIPSNO` as the `Key Variable`. Select `Tools > Data Export > Ascii`. In the dialog, specify `south.dbf` as the input file and `south60.txt` as the output file. Select `FIPSNO`, `HR60`, `RD60`, `PS60`, `MA60`, `DV60` and `UE60`.[1]

---

[1]This replicates the analysis in Baller et al. (2001), although the spatial weights used

Figure 6.1: Cumulative first and second order rook weights

Next, click on `Export` and `Done` to finish. Check the contents of the `south.txt` file and make sure it shows `1412,7` on the first line and the variable names on the second line. The data, record by record, and comma-delimited make up the remainder. Copy the file to the R working directory if it is not there already.

Repeat the same procedure to create a file `south90.txt` with the variables `FIPSNO`, `HR90`, `RD90`, `PS90`, `MA90`, `DV90` and `UE90`.

Also create two spatial weights files. One is a simple rook contiguity (`southrk.gal`), the other one that uses a cumulative first and second order rook criterion. For the latter, in *GeoDa*, select `Tools > Weights > Create`, enter `south.shp` as the input file and `southrk12.gal` as the output file and `FIPSNO` as the ID. Select the `rook` criterion, change the `order of contiguity` to 2 and check the box to `Include all the lower orders`, as shown in Figure 6.1. For the first order rook weights, proceed in the same way, but leave the option to `rook`.

Copy all the files to the R working directory. First, create a separate data frame for the `south60` and `south90` data and make sure the variables are entered properly, as in:

```
> south60 <- read.csv("south60.txt",header=TRUE,skip=1)
> summary(south60)
     FIPSNO              HR60                RD60
 Min.   : 1001   Min.   : 0.000   Min.    :-1.85781
 1st Qu.:13284   1st Qu.: 3.213   1st Qu.: 0.04904
 Median :37088   Median : 6.245   Median : 0.66042
 Mean   :31936   Mean   : 7.292   Mean    : 0.68974
 3rd Qu.:48158   3rd Qu.: 9.956   3rd Qu.: 1.32815
 Max.   :54109   Max.   :92.937   Max.    : 3.45354
 ...
```

here are slightly different. Check the article and the code book for the `nat.shp` sample data set for a detailed explanation of the variables and theoretical background.

81

and

```
> south90 <- read.csv("south90.txt",header=TRUE,skip=1)
> summary(south90)
     FIPSNO            HR90              RD90
 Min.   : 1001   Min.   : 0.000   Min.    :-2.2006
 1st Qu.:13284   1st Qu.: 4.805   1st Qu.:-0.1021
 Median :37088   Median : 8.221   Median : 0.3720
 Mean   :31936   Mean   : 9.549   Mean    : 0.5536
 3rd Qu.:48158   3rd Qu.:13.038   3rd Qu.: 1.1203
 Max.   :54109   Max.   :64.261   Max.    : 5.5831
 . . .
```

Next, with spdep active, turn both gal text files into a neighbor object by means of the read.gal function, as in:

```
> srk12 <- read.gal("southrk12.gal",override.id=TRUE)
> summary(srk12)
Neighbour list object:
Number of regions: 1412
Number of nonzero links: 23768
Percentage nonzero weights: 1.192129
Average number of links: 16.83286
Link number distribution:

  2   3   4   5   6   7   8
  1   4   4   8   9  11  31
  . . .
1 least connected region:
54029 with 2 links
1 most connected region:
51015 with 29 links
```

and finally into a listw object as in:

```
> srk12w <- nb2listw(srk12)
> summary(srk12w)
Characteristics of weights list object:
Neighbour list object:
Number of regions: 1412
Number of nonzero links: 23768
```

```
Percentage nonzero weights: 1.192129
Average number of links: 16.83286
Link number distribution:

  2   3   4   5   6   7   8
  1   4   4   8   9  11  31
  . . .
1 least connected region:
54029 with 2 links
1 most connected region:
51015 with 29 links

Weights style: W
Weights constants summary:
     n      nn     S0       S1        S2
W 1412 1993744 1412 180.5536 5703.863
```

Proceed in the same way for `southrk.gal`:

```
> srk <- read.gal("southrk.gal",override.id=TRUE)
> srkw <- nb2listw(srk)
> summary(srkw)
Characteristics of weights list object:
Neighbour list object:
Number of regions: 1412
Number of nonzero links: 7700
Percentage nonzero weights: 0.3862081
Average number of links: 5.453258
Link number distribution:

  1   2   3   4   5   6   7   8
 16  32  65 187 378 435 230  56


 ....

Weights style: W
Weights constants summary:
     n      nn     S0       S1        S2
W 1412 1993744 1412 552.1101 5767.949
```

We are now ready to start the analysis.

6.3 Baseline OLS Regression

We start by creating an `lm` *object* that will need to be passed to the specialized functions that carry out the spatial autocorrelation tests. We use the familiar `lm` function. We start with the analysis for 1960 with `HR60` as dependent variable and explanatory variables `RD60 + PS60 + MA60 + DV60 + UE60`. First, we `attach` the `south60` data frame so that we can easily refer to the variables:

```
> attach(south60)
> sols60 <- lm(HR60 ~ RD60 + PS60 + MA60 + DV60 + UE60)
> summary(sols60)

Call:
lm(formula = HR60 ~ RD60 + PS60 + MA60 + DV60 + UE60)

Residuals:
     Min       1Q   Median       3Q      Max
-13.5265  -3.6178  -0.8159   2.4218  88.4821

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 13.21547    1.12457  11.752  < 2e-16 ***
RD60         1.76448    0.19824   8.901  < 2e-16 ***
PS60         0.29930    0.21426   1.397    0.163
MA60        -0.27521    0.03806  -7.230 7.89e-13 ***
DV60         1.17945    0.24352   4.843 1.42e-06 ***
UE60        -0.29186    0.07117  -4.101 4.35e-05 ***
---
Signif. codes:  0 *** 0.001 ** 0.01 * 0.05 . 0.1   1

Residual standard error: 6.09 on 1406 degrees of freedom
Multiple R-Squared: 0.1037,Adjusted R-squared: 0.1005
F-statistic: 32.52 on 5 and 1406 DF,  p-value: < 2.2e-16
```

The values are identical to those listed in the *GeoDa Workbook* (Figure 24.3), as they should be.

We next create a similar object for the 1990 analysis, using the same variables (with a `90` posfix). First detach the `south60` data frame and attach the `south90` data frame.

84

```
> detach(south60)
> attach(south90)
> sols90 <- lm(HR90 ~ RD90 + PS90 + MA90 + DV90 + UE90)
> summary(sols90)

Call:
lm(formula = HR90 ~ RD90 + PS90 + MA90 + DV90 + UE90)

Residuals:
     Min       1Q    Median       3Q      Max
-19.5829  -3.5137   -0.7474   2.5908  41.8327

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  8.96250    1.78133   5.031 5.50e-07 ***
RD90         4.58778    0.21457  21.381  < 2e-16 ***
PS90         1.95589    0.20540   9.522  < 2e-16 ***
MA90        -0.04948    0.04890  -1.012    0.312
DV90         0.46159    0.11517   4.008 6.45e-05 ***
UE90        -0.52440    0.07003  -7.489 1.22e-13 ***
---
Signif. codes:  0 *** 0.001 ** 0.01 * 0.05 . 0.1   1

Residual standard error: 5.861 on 1406 degrees of freedom
Multiple R-Squared: 0.3092,Adjusted R-squared: 0.3067
F-statistic: 125.8 on 5 and 1406 DF,  p-value: < 2.2e-16
```

These results are identical to what is listed in Figure 25.2 of the *GeoDa Workbook*. Relative to 1960, the overall fit is better in 1990. The significances and signs are generally the same, except for PS not being significant in 1960 while significant in 1990, and the reverse for MA.

6.4   Moran's I for Regression Residuals

In order to carry out Moran's I test on the residuals in these regressions, we need to pass the regression object and a spatial weights object (listw) to the lm.morantest function (check the help file for details). The print command provides the results.

We start with the analysis for 1960 using the cumulative weights:

```
> sols60.moran <- lm.morantest(sols60,srk12w)
> sols60.moran

Global Moran's I for regression residuals

data:
model: lm(formula = HR60 ~ RD60 + PS60 + MA60 + DV60 + UE60)
weights: srk12w

Moran I statistic standard deviate = 14.8749, p-value < 2.2e-16
alternative hypothesis: greater
sample estimates:
Observed Moran's I        Expectation          Variance
      1.365683e-01       -2.215366e-03       8.705017e-05
```

Note that the default setting for this statistic is to compute the p-value for a one sided test. To get a two-sided test, the `alternative` argument must be specified explicitly, as in:

```
> lm.morantest(sols60,srk12w,alternative="two.sided")

Global Moran's I for regression residuals

data:
model: lm(formula = HR60 ~ RD60 + PS60 + MA60 + DV60 + UE60)
weights: srk12w

Moran I statistic standard deviate = 14.8749, p-value < 2.2e-16
alternative hypothesis: two.sided
sample estimates:
Observed Moran's I        Expectation          Variance
      1.365683e-01       -2.215366e-03       8.705017e-05
```

In this example, it does not make a difference, due to the large sample size and the extreme significance of the test statistic. However, in small samples, inference will be affected when the wrong option is used.

For the simple first order rook contiguity weights, the results are very similar:

```
> lm.morantest(sols60,srkw,alternative="two.sided")
```

```
Global Moran's I for regression residuals

data:
model: lm(formula = HR60 ~ RD60 + PS60 + MA60 + DV60 + UE60)
weights: srkw

Moran I statistic standard deviate = 8.3495, p-value < 2.2e-16
alternative hypothesis: two.sided
sample estimates:
Observed Moran's I          Expectation              Variance
      0.1356298451       -0.0024963269          0.0002736731
```

The results show a Moran's I statistic of respectively 0.137 and 0.136, which are highly significant and reject the null hypothesis of uncorrelated error terms.

For 1990, similar results are obtained, with slightly smaller values for the test statistic, 0.122 using the simple rook, and 0.009 using the cumulative rook:

```
> lm.morantest(sols90,srkw,alternative="two.sided")

Global Moran's I for regression residuals

data:
model: lm(formula = HR90 ~ RD90 + PS90 + MA90 + DV90 + UE90)
weights: srkw

Moran I statistic standard deviate = 7.5134, p-value = 5.76e-14
alternative hypothesis: two.sided
sample estimates:
Observed Moran's I          Expectation              Variance
      0.1218894764       -0.0023982682          0.0002736415
```

and

```
> lm.morantest(sols90,srk12w,alternative="two.sided")

Global Moran's I for regression residuals

data:
model: lm(formula = HR90 ~ RD90 + PS90 + MA90 + DV90 + UE90)
```

```
weights: srk12w

Moran I statistic standard deviate = 9.8644, p-value < 2.2e-16
alternative hypothesis: two.sided
sample estimates:
Observed Moran's I        Expectation            Variance
      8.993101e-02       -2.133033e-03         8.710495e-05
```

6.5 Lagrange Multiplier Test Statistics

The Morans I test statistic has high power against a range of spatial alternatives. However, it does not provide much help in terms of which alternative model would be most appropriate. The Lagrange Multiplier test statistics do allow a distinction between spatial error models and spatial lag models.

Both tests, as well as their robust forms are included in the `lm.LMtests` function. Again, a regression object and a spatial `listw` object must be passed as arguments. In addition, the tests must be specified as a character vector (the default is only `LMerror`), using the `c( )` operator (concatenate), as illustrated below.

First consider the 1960 results, for both spatial weights.

```
> lm.LMtests(sols60,srk12w,
+ test=c("LMerr","RLMerr","LMlag","RLMlag"))

Lagrange multiplier diagnostics for spatial dependence

data:
model: lm(formula = HR60 ~ RD60 + PS60 + MA60 + DV60 + UE60)
weights: srk12w
LMerr = 205.9506, df = 1, p-value < 2.2e-16


Lagrange multiplier diagnostics for spatial dependence

data:
model: lm(formula = HR60 ~ RD60 + PS60 + MA60 + DV60 + UE60)
weights: srk12w
RLMerr = 1.8681, df = 1, p-value = 0.1717


Lagrange multiplier diagnostics for spatial dependence
```

```
data:
model: lm(formula = HR60 ~ RD60 + PS60 + MA60 + DV60 + UE60)
weights: srk12w
LMlag = 222.528, df = 1, p-value < 2.2e-16


Lagrange multiplier diagnostics for spatial dependence

data:
model: lm(formula = HR60 ~ RD60 + PS60 + MA60 + DV60 + UE60)
weights: srk12w
RLMlag = 18.4456, df = 1, p-value = 1.748e-05


> lm.LMtests(sols60,srkw,
+ test=c("LMerr","RLMerr","LMlag","RLMlag"))


Lagrange multiplier diagnostics for spatial dependence

data:
model: lm(formula = HR60 ~ RD60 + PS60 + MA60 + DV60 + UE60)
weights: srkw
LMerr = 66.4285, df = 1, p-value = 3.331e-16


Lagrange multiplier diagnostics for spatial dependence

data:
model: lm(formula = HR60 ~ RD60 + PS60 + MA60 + DV60 + UE60)
weights: srkw
RLMerr = 5.1188, df = 1, p-value = 0.02367


Lagrange multiplier diagnostics for spatial dependence

data:
model: lm(formula = HR60 ~ RD60 + PS60 + MA60 + DV60 + UE60)
weights: srkw
LMlag = 80.3219, df = 1, p-value < 2.2e-16


Lagrange multiplier diagnostics for spatial dependence

data:
model: lm(formula = HR60 ~ RD60 + PS60 + MA60 + DV60 + UE60)
```

```
weights: srkw
RLMlag = 19.0122, df = 1, p-value = 1.299e-05
```

The results are listed in order with `LMerr` and `RLMerr` first, followed by `LMlag` and `RLMlag`. The values for the `srk12w` weights correspond to the results shown in Figure 24.4 of the *GeoDa Workbook*. Both `LMerr` and `LMlag` are highly significant, with a slight edge for the latter. In a specification search, this is an inconclusive results, necessitating the consideration of the robust forms of the tests.

These provide a clear preference in favor of the lag alternative. While `RLMlag` is still strongly significant, for both weights, there is much weaker evidence given by `RLMerr`. For the `srkw` weights it is still marginally significant, but much less so than the robust lag test. However, for the `srk12w` weights it is no longer significant (p = 0.17), clearly pointing to the lag model as the proper alternative.

Interestingly, the results are the opposite for 1990:

```
> lm.LMtests(sols90,srk12w,
+ test=c("LMerr","RLMerr","LMlag","RLMlag"))


Lagrange multiplier diagnostics for spatial dependence

data:
model: lm(formula = HR90 ~ RD90 + PS90 + MA90 + DV90 + UE90)
weights: srk12w
LMerr = 89.3063, df = 1, p-value < 2.2e-16


Lagrange multiplier diagnostics for spatial dependence

data:
model: lm(formula = HR90 ~ RD90 + PS90 + MA90 + DV90 + UE90)
weights: srk12w
RLMerr = 22.3826, df = 1, p-value = 2.234e-06


Lagrange multiplier diagnostics for spatial dependence

data:
model: lm(formula = HR90 ~ RD90 + PS90 + MA90 + DV90 + UE90)
weights: srk12w
LMlag = 71.6978, df = 1, p-value < 2.2e-16
```

```
Lagrange multiplier diagnostics for spatial dependence

data:
model: lm(formula = HR90 ~ RD90 + PS90 + MA90 + DV90 + UE90)
weights: srk12w
RLMlag = 4.774, df = 1, p-value = 0.02889

> lm.LMtests(sols90,srkw,
+ test=c("LMerr","RLMerr","LMlag","RLMlag"))

Lagrange multiplier diagnostics for spatial dependence

data:
model: lm(formula = HR90 ~ RD90 + PS90 + MA90 + DV90 + UE90)
weights: srkw
LMerr = 53.6508, df = 1, p-value = 2.395e-13

Lagrange multiplier diagnostics for spatial dependence

data:
model: lm(formula = HR90 ~ RD90 + PS90 + MA90 + DV90 + UE90)
weights: srkw
RLMerr = 5.5171, df = 1, p-value = 0.01883

Lagrange multiplier diagnostics for spatial dependence

data:
model: lm(formula = HR90 ~ RD90 + PS90 + MA90 + DV90 + UE90)
weights: srkw
LMlag = 50.5794, df = 1, p-value = 1.144e-12

Lagrange multiplier diagnostics for spatial dependence

data:
model: lm(formula = HR90 ~ RD90 + PS90 + MA90 + DV90 + UE90)
weights: srkw

RLMlag = 2.4457, df = 1, p-value = 0.1178
```

Again, both `LMerr` and `LMlag` are highly significant for both weights. However, the robust statistics point to the error model as the proper alternative. For `srk12w`, `RLMlag` is only weakly significant (p = 0.03), but `RLMerr` is still very strongly significant. The evidence is clearer for `srkw`, where `RLMlag` reaches a p-value of 0.12, whereas `RLMerr` is significant with p=0.02.[2]

With this information in hand, we can select the spatial lag model as the alternative for 1960 and the spatial error model as the alternative for 1990. Given the magnitude of the test statistics for the different weights, we opt for `srkw12` for the lag model in 1960 and `srkw` for the error model in 1990. Maximum Likelihood estimation is covered in the next chapter.

6.6   Practice

Run the diagnostics for the same specification in years 1970 and 1980 and assess whether a spatial lag or error model is the proper model. Alternatively, use the POLICE sample data set and rerun the regression from Kelejian and Robinson (1992), using `POLICE` as the dependent variable, and `TAX`, `INC`, `CRIME`, `UNEMP`, `OWN`, `COLLEGE`, `WHITE` and `COMMUTE` as explanatory variables (see Kelejian and Robinson 1992, pp. 323–324). Does your analysis support the conclusions found in that article. Experiment with different spatial weights to assess the sensitivity of your findings.

---

[2]These results are also shown in Figure 25.3 of the *GeoDa Workbook*.

Exercise 7

Maximum Likelihood Estimation of Spatial Regression Models

7.1 Objectives

This exercise illustrates how to estimate spatial regression models using the maximum likelihood method. Both spatial lag and spatial error models are included. The example continues the analysis of Southern US county homicides from the previous exercise and matches the results in Chapters 24 and 25 of the *GeoDa Workbook* (Anselin 2005).

7.2 Preliminaries

Make sure to have the `spdep` library active before starting the exerices, or else invoke `library(spdep)`. Also, you will need the data frames and spatial weights for the 1960 and 1990 homicides used in Exercise 6. Refer to Section 6.2 to recreate these if necessary.

7.3 ML Estimation of the Spatial Lag Model

Maximum Likelihood (ML) estimation of the spatial lag model is carried out with the `lagsarlm( )` function. The required arguments are a regression formula, a data set and a listw spatial weights object. The default method uses Ords eigenvalue decomposition of the spatial weights matrix (Ord 1975).

There is also a sparse weights method available which may be necessary for large(r) data sets. However, the inference provided by the sparse method is limited (for details, see the help file).

In our example, we will use the `srk12w` weights and the `south60` data frame, with the same specification as for the OLS regression. The dependent variable is `HR60`, and the explanatory variables are `RD60 + PS60 + MA60 + DV60 + UE60`.

We create a new object as the output of the estimation procedure. This will allow us to extract specific results, such as coefficients, residuals and predicted values.

```
> lag60 <- lagsarlm(HR60 ~ RD60 + PS60 + MA60 + DV60 + UE60,
+data=south60,srk12w)
> summary(lag60)

Call:lagsarlm(formula = HR60 ~ RD60 + PS60 + MA60 + DV60 + UE60,
        data = south60, listw = srk12w)

Residuals:
      Min        1Q    Median        3Q       Max
-14.97132  -3.25129  -0.79033   2.22485  86.27312


Type: lag
Coefficients: (asymptotic standard errors)
            Estimate Std. Error z value  Pr(>|z|)
(Intercept)  6.574961   1.172724  5.6066 2.064e-08
RD60         1.100473   0.196339  5.6050 2.083e-08
PS60         0.037912   0.202678  0.1871   0.85162
MA60        -0.175256   0.036712 -4.7738 1.808e-06
DV60         0.935208   0.230386  4.0593 4.922e-05
UE60        -0.132660   0.067353 -1.9696   0.04888


Rho: 0.53289 LR test value: 125.07 p-value: < 2.22e-16
Asymptotic standard error: 0.045668 z-value: 11.669 p-value: < 2.22e-16
Wald statistic: 136.16 p-value: < 2.22e-16

Log likelihood: -4488.967 for lag model
ML residual variance (sigma squared): 33.045, (sigma: 5.7485)
Number of observations: 1412
Number of parameters estimated: 8
```

```
AIC: 8993.9, (AIC for lm: 9117)
LM test for residual autocorrelation
test value: 4.3536 p-value: 0.03693
```

The results match those reported in Figure 24.9 of the *GeoDa Workbook* even though *GeoDa* uses a different numerical approach to maximize the likelihood function. Some caution is needed with the eigenvalue approach used in the `sarlaglm` function, since it may be numerically unstable for large data sets ($> 1,000$). In the current example, that does not seem to be an issue. Note that compared to the OLS results, `PS` is still not significant, but now `UE` has become only marginally significant. The autoregressive coefficient of `0.533` is high in magnitude and highly significant.

The output also contains a LM test statistic for remaining spatial error autocorrelation. The value of `4.35` is marginally significant ($p = 0.037$) suggesting remaining issues with the specification of the spatial dependence.

Note how different the output is when you specify `Matrix` as the estimation method instead of the default `eigen`:

```
> lag60s <- lagsarlm(HR60 ~ RD60 + PS60 + MA60 + DV60 + UE60,
+ data=south60,srk12w,method="Matrix")
> summary(lag60s)

Call:lagsarlm(formula = HR60 ~ RD60 + PS60 + MA60 + DV60 + UE60,
    data = south60, listw = srk12w, method = "Matrix")

Residuals:
      Min        1Q    Median        3Q       Max
-14.97132  -3.25129  -0.79033   2.22485  86.27312


Type: lag
Coefficients: (log likelihood/likelihood ratio)
              Estimate Log likelihood LR statistic  Pr(>|z|)
(Intercept)  6.5750e+00             NA           NA        NA
RD60         1.1005e+00     -4.5049e+03      31.8236 1.688e-08
PS60         3.7912e-02     -4.4890e+03       0.0348   0.85211
MA60        -1.7526e-01     -4.5002e+03      22.5635 2.033e-06
DV60         9.3521e-01     -4.4971e+03      16.3583 5.243e-05
UE60        -1.3266e-01     -4.4908e+03       3.7600   0.05249

Rho: 0.53289 LR test value: 125.07 p-value: < 2.22e-16
```

```
Log likelihood: -4488.967 for lag model
ML residual variance (sigma squared): 33.045, (sigma: 5.7485)
Number of observations: 1412
Number of parameters estimated: 8
AIC: 8993.9, (AIC for lm: 9117)
```

While the coefficient estimates are identical, there are quite some differences in terms of the inference. First of all, there is no asymptotic variance for any of the coefficients. Instead, significance is based on a Likelihood Ratio test (LR). For the autoregressive coefficient, this is the LR test for the full model. This shows the same value of 125.07 as with the eigenvalue method. There is no inference for the constant term. For the regression coefficients, the value reported in the column customarily reserved for the asymptotic standard error is actually the log likelihood of a model where this variable has been dropped from the specification. For example, for RD60, this value is -4504.9. The next item, customarily reserved for the t-value, is the LR test statistic from comparing the log likelihood of the unrestricted model (with the variable) to that of the restricted model (without the variable). Formally, $LR = -2(L_U - L_R)$. In our example, the negative difference between -4488.967 and -4504.879 is 15.912, and twice this value (the LR test statistic) is 31.824, the value listed in the third column next to RD60. The fourth column shows the p-value for the LR test.

These individual log-likelihoods can also be extracted explicitly from the fitted object, as the list item LLs, using the familiar $ notation to extract an element from a list:

```
> lag60s$LLs
[[1]]
'log Lik.' -4504.879 (df=7)
[[2]]
'log Lik.' -4488.985 (df=7)
[[3]]
'log Lik.' -4500.249 (df=7)
[[4]]
'log Lik.' -4497.146 (df=7)
[[5]]
'log Lik.' -4490.847 (df=7)
```

Other items of interest are the residuals, e.g., lag60$residuals, and the predicted (fitted) values, e.g., lag60$fitted.values. Note that lagsarlm

uses a different approach from *GeoDa* to compute the predicted values. In *GeoDa* the latter are obtained from the reduced form, as $\hat{y} = (I - \hat{\rho}W)^{-1}X\hat{\beta}$, whereas `lagsarlm` uses $\hat{y} = \hat{\rho}Wy + X\hat{\beta}$. See Section 24.4 of the *GeoDa Workbook* for an in-depth discussion of this issue and some illustrations. See also the `lagsarlm` help file for further technical details on its implementation.

7.4   ML Estimation of the Spatial Error Model

ML estimation of the spatial error model is similar to the lag procedure and implemented in the `errorsarlm( )` function. Again, the formula, data set and a listw spatial weights object must be specified at a minimum. The default estimation procedure is the eigenvalue decomposition.

To illustrate this, consider the estimation of the homicide specification for 1990, using `south90` as the data set, `srkw` as the weights matrix, and the specification with `HR90` as the dependent variable, and `RD90 + PS90 + MA90 + DV90 + UE90` as the explanatory variables.

```
> err90 <- errorsarlm(HR90 ~ RD90 + PS90 + MA90 + DV90 + UE90,
+ data=south90,srkw)
> summary(err90)

Call:errorsarlm(formula = HR90 ~ RD90 + PS90 + MA90 + DV90 + UE90,
     data = south90, listw = srkw)

Residuals:
      Min        1Q    Median        3Q       Max
-18.17590  -3.50131  -0.63837   2.43642  41.70782


Type: error
Coefficients: (asymptotic standard errors)
             Estimate Std. Error z value  Pr(>|z|)
(Intercept)  6.693514   1.958045   3.4185 0.0006297
RD90         4.407396   0.237668  18.5443 < 2.2e-16
PS90         1.766328   0.225652   7.8276 4.885e-15
MA90        -0.016640   0.052990  -0.3140 0.7535090
DV90         0.499146   0.124912   3.9960 6.443e-05
UE90        -0.387841   0.078478  -4.9420 7.731e-07


Lambda: 0.29161 LR test value: 52.11 p-value: 5.2491e-13
Asymptotic standard error: 0.037275 z-value: 7.8231 p-value: 5.107e-15
```

```
Wald statistic: 61.201 p-value: 5.107e-15

Log likelihood: -4471.317 for error model
ML residual variance (sigma squared): 32.407, (sigma: 5.6927)
Number of observations: 1412
Number of parameters estimated: 8
AIC: 8958.6, (AIC for lm: 9008.7)
```

The results are identical to those listed in Figure 25.6 of the *GeoDa Workbook*, even though a different numerical optimization approach is used. The results are similar to those obtained with OLS (`MA` remains insignificant). Again, when specifying the method `Matrix` the inference differs but not in the same way as for the lag case.

For example, consider the results below:

```
> err90s <- errorsarlm(HR90 ~ RD90 + PS90 + MA90 + DV90 + UE90,
+ data=south90,srkw,method="Matrix")
> summary(err90s)

Call:errorsarlm(formula = HR90 ~ RD90 + PS90 + MA90 + DV90 + UE90,
    data = south90, listw = srkw, method = "Matrix")

Residuals:
      Min        1Q    Median        3Q       Max
-18.17590  -3.50131  -0.63837   2.43642  41.70782

Type: error
Coefficients: (asymptotic standard errors)
            Estimate Std. Error z value  Pr(>|z|)
(Intercept)  6.693514   1.958045  3.4185 0.0006297
RD90         4.407397   0.237668 18.5443 < 2.2e-16
PS90         1.766328   0.225652  7.8277 4.885e-15
MA90        -0.016640   0.052990 -0.3140 0.7535090
DV90         0.499146   0.124912  3.9960 6.443e-05
UE90        -0.387841   0.078478 -4.9420 7.731e-07

Lambda: 0.29161 LR test value: 52.11 p-value: 5.2491e-13

Log likelihood: -4471.317 for error model
ML residual variance (sigma squared): 32.407, (sigma: 5.6927)
```

```
Number of observations: 1412
Number of parameters estimated: 8
AIC: 8958.6, (AIC for lm: 9008.7)
```

There is no asymptotic error variance for the autoregressive coefficient. As for the lag, a LR test is used instead. The inference for the regression coefficients is the same as for the eigenvector approach. This is because, unlike what holds for the lag model, the asymptotic variance matrix in the error model is block diagonal between the spatial parameter and the other parameters of the model. The asymptotic variance for the latter is obtained by plugging the estimate for the spatial autoregressive coefficient into the familiar FGLS expression. In the lag model, there is no such simple block diagonal shape.

Specific components of the fitted model can be extracted in the usual fashion. For example, residuals are obtained as `err90$residuals` and predicted values as `err90$fitted.values`. See the help file for details.

7.5   Spatial Durbin Model

An important specification test in the spatial error model is a test on the so-called spatial common factor hypothesis. This exploits the property that a spatial error model can also be specified in spatial lag form, with the spatially lagged explanatory variables included, but with constraints on the parameters (the common factor constraints). The spatial lag form of the error model is also referred to as the spatial Durbin specification. In `spdep`, this is implemented by specifying the `mixed` option for type in the spatial lag estimation (the spatially lagged explanatory variables need not be specified).

To assess whether the common factor constraints are satisfied for the spatial error model of `HR90`, we run the same specification as a lag model, with the `type="mixed"` option.

```
> lag90durbin <- lagsarlm(HR90 ~ RD90 + PS90 + MA90 + DV90 + UE90,
+ data=south90,srkw,type="mixed")
> summary(lag90durbin)

Call:lagsarlm(formula = HR90 ~ RD90 + PS90 + MA90 + DV90 + UE90,
    data = south90, listw = srkw, type = "mixed")

Residuals:
     Min        1Q    Median        3Q       Max
```

```
-18.26433  -3.30003  -0.69833   2.43177  39.77795


Type: mixed
Coefficients: (asymptotic standard errors)
             Estimate Std. Error z value  Pr(>|z|)
(Intercept) 10.861374   2.571244   4.2242 2.398e-05
RD90         4.176400   0.283976 14.7069 < 2.2e-16
PS90         1.373438   0.258171   5.3199 1.038e-07
MA90         0.035280   0.059578   0.5922 0.5537373
DV90         0.561657   0.139475   4.0269 5.651e-05
UE90        -0.137848   0.093086  -1.4809 0.1386421
lag.RD90    -0.441362   0.431248  -1.0235 0.3060940
lag.PS90     0.536691   0.370724   1.4477 0.1477053
lag.MA90    -0.152049   0.090370  -1.6825 0.0924669
lag.DV90    -0.224587   0.201777  -1.1130 0.2656886
lag.UE90    -0.482028   0.127317  -3.7860 0.0001531


Rho: 0.2547 LR test value: 41.96 p-value: 9.318e-11
Asymptotic standard error: 0.037956 z-value: 6.7105 p-value: 1.9398e-11
Wald statistic: 45.031 p-value: 1.9398e-11


Log likelihood: -4453.823 for mixed model
ML residual variance (sigma squared): 31.747, (sigma: 5.6344)
Number of observations: 1412
Number of parameters estimated: 13
AIC: 8933.6, (AIC for lm: 8973.6)
LM test for residual autocorrelation
test value: 12.986 p-value: 0.00031377
```

The output lists the estimates, standard errors, etc. for each explanatory variable as well as its spatially lagged counterpart (e.g., `lag.RD90`). Several of the coefficient estimates for the latter are not significant. Note also how the spatial autoregressive coefficient of `0.2547` is slightly different from the value obtained with the pure error model (`0.2916`). Also, there is evidence of remaining spatial error autocorrelation, as indicated by the signficant LM test statistic.

The main purpose of this exercise is to test whether the product of this spatial autoregressive coefficient with each of the regression coefficients equals negative the coefficient of the matching spatially lagged explanatory variable. There are a total of 5 such constraints, which gives the degrees of

freedom for the test on the common factor constraints.

This is implemented through the `LR.sarlm` function, which allows a likelihood ratio test between two *nested* models (i.e., one must be found from the other by setting coefficients equal to zero). In this case, the nested models are the unconstrained spatial Durbin model (`lag90durbin`) and the constrained spatial error model (`err90`). This yields:

```
> LR.sarlm(lag90durbin,err90)

Likelihood ratio for spatial linear models

data:
Likelihood ratio = 34.9882, df = 5, p-value = 1.513e-06
sample estimates:
Log likelihood of lag90durbin       Log likelihood of err90
                  -4453.823                      -4471.317
```

The constraints are strongly rejected. In other words, the spatial autoregressive error specification is internally inconsistent. This can be due to a poor choice of the spatial weights, or to a misspecification of the error process (e.g., when the true error process is not SAR). The upshot of this evidence is that the model specification should be reconsidered.

7.6   Practice

Experiment with different spatial weights to obtain a more satisfactory spatial regression model for `HR90`. Alternatively, explore the spatial regression estimation in one of the sample data sets, e.g., replicating the POLICE example, or the Baltimore or Boston house price examples.

Exercise 8

Discrete Spatial Heterogeneity

8.1   Objectives

While `spdep` does not contain specific functionality to address spatial heterogeneity, it is possible to get some useful results by manipulating the regression format and extracting information from the fitted spatial regression object. This exercise illustrates the analysis of discrete spatial heterogeneity in the form of spatial analysis of variance (spatial ANOVA) and tests for the presence of spatial regimes.

8.2   Preliminaries

As a running example, we will continue the analysis of homicide rates in US counties, but now consider the whole country. Follow the instructions given in Section 6.2 to create a data frame from the sample data set `nat.shp`, selecting the variables `FIPSNO`, `SOUTH`, `HR60`, `RD60`, `PS60`, `MA60`, `DV60` and `UE60`. Turn this into the data frame `nat60`.

   In addition, create a neighbor and listw object from the queen contiguity file `natqueen.gal` using the `read.gal` and `nb2listw` functions. Call the final listw object `natq`.

   Make sure everything is correct by comparing your data summaries to the results below:

```
> summary(nat60)
     FIPSNO            SOUTH             HR60
 Min.   : 1001    Min.   :0.0000    Min.   : 0.000
 1st Qu.:19037    1st Qu.:0.0000    1st Qu.: 0.000
```

```
Median :29197    Median :0.0000    Median : 2.783
Mean   :30592    Mean   :0.4577    Mean   : 4.504
3rd Qu.:45083    3rd Qu.:1.0000    3rd Qu.: 6.885
Max.   :56045    Max.   :1.0000    Max.   :92.937
 . . .
```

and

```
> summary(natq)
Characteristics of weights list object:
Neighbour list object:
Number of regions: 3085
Number of nonzero links: 18168
Percentage nonzero weights: 0.1908960
Average number of links: 5.889141
Link number distribution:

   1    2    3    4    5    6    7    8    9
  24   36   91  281  620 1037  704

 . . .

1 most connected region:
49037 with 14 links

Weights style: W
Weights constants summary:
    n       nn   S0       S1       S2
W 3085 9517225 3085 1104.156 12568.36
```

8.3   Spatial ANOVA

Spatial analysis of variance can be implemented as a simple dummy variable regression. In our example, HR90 is the dependent variable, with the dummy variable SOUTH as the explanatory variable. The estimate of the dummy indicates the difference between the *regime* SOUTH = 1 and the overall mean.

8.3.1   OLS Estimation

We first implement this using the standard lm function and create a regression object to test for spatial error autocorrelation.

Before you start, make sure to `attach` the `nat60` data frame. The results show a highly significant effect of the South.

```
> attach(nat60)
> hr60an <- lm(HR60 ~ SOUTH)
> summary(hr60an)

Call:
lm(formula = HR60 ~ SOUTH)

Residuals:
   Min     1Q Median     3Q    Max
-7.292 -2.151 -1.052  1.467 85.645

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)   2.1510     0.1231   17.47   <2e-16 ***
SOUTH         5.1412     0.1820   28.25   <2e-16 ***
---
Signif. codes:  0 *** 0.001 ** 0.01 * 0.05 . 0.1   1

Residual standard error: 5.036 on 3083 degrees of freedom
Multiple R-Squared: 0.2056,Adjusted R-squared: 0.2053
F-statistic: 797.9 on 1 and 3083 DF,  p-value: < 2.2e-16
```

We next check the LM statistics for spatial autocorrelation:

```
> lm.LMtests(hr60an,natq,
+ test=c("LMerr","RLMerr","LMlag","RLMlag"))

Lagrange multiplier diagnostics for spatial dependence

data:
model: lm(formula = HR60 ~ SOUTH)
weights: natq
LMerr = 332.4194, df = 1, p-value < 2.2e-16

Lagrange multiplier diagnostics for spatial dependence

data:
```

```
model: lm(formula = HR60 ~ SOUTH)
weights: natq
RLMerr = 13.1194, df = 1, p-value = 0.0002923


Lagrange multiplier diagnostics for spatial dependence

data:
model: lm(formula = HR60 ~ SOUTH)
weights: natq
LMlag = 359.7856, df = 1, p-value < 2.2e-16


Lagrange multiplier diagnostics for spatial dependence

data:
model: lm(formula = HR60 ~ SOUTH)
weights: natq
RLMlag = 40.4856, df = 1, p-value = 1.981e-10
```

The results show very strong evidence of spatial autocorrelation, with a slight edge in favor of the lag alternative. We will estimate both and assess the effect on the magnitude and significance of the coefficient of the SOUTH dummy.

8.3.2 Spatial Dummy Variable Regression

For completeness sake, we check both lag and error alternatives. The spatial lag version of the ANOVA model yields:

```
> hr60anlag <- lagsarlm(HR60 ~ SOUTH,data=nat60,
+ natq)
> summary(hr60anlag)

Call:lagsarlm(formula = HR60 ~ SOUTH, data = nat60, listw = natq)

Residuals:
      Min        1Q    Median        3Q       Max
-12.99088  -2.11958  -0.93541   1.37336  86.54842


Type: lag
Coefficients: (asymptotic standard errors)
           Estimate Std. Error z value  Pr(>|z|)
```

```
(Intercept)   1.29761     0.12815   10.125 < 2.2e-16
SOUTH         3.16222     0.20923   15.114 < 2.2e-16

Rho: 0.38603 LR test value: 261.92 p-value: < 2.22e-16
Asymptotic standard error: 0.024275 z-value: 15.902 p-value: < 2.22e-16
Wald statistic: 252.88 p-value: < 2.22e-16

Log likelihood: -9232.926 for lag model
ML residual variance (sigma squared): 22.626, (sigma: 4.7566)
Number of observations: 3085
Number of parameters estimated: 4
AIC: 18474, (AIC for lm: 18734)
LM test for residual autocorrelation
test value: 151.76 p-value: < 2.22e-16
```

Compared to the OLS results, the dummy remains strongly significant, even though its magnitude decreases considerably after the spatial lag *filter* is implemented. A cursory look at the diagnostics shows that considerably spatial autocorrelation remains, which is not surprising, given the simplicity of the specification.

The spatial error version of the ANOVA model yields:

```
> hr60anerr <- errorsarlm(HR60 ~ SOUTH,data=nat60,
+ natq)
> summary(hr60anerr)

Call:errorsarlm(formula = HR60 ~ SOUTH, data = nat60, listw = natq)

Residuals:
      Min        1Q     Median        3Q       Max
-12.75378  -2.15718  -0.95972   1.42182  86.62697

Type: error
Coefficients: (asymptotic standard errors)
            Estimate Std. Error z value  Pr(>|z|)
(Intercept) 2.24729     0.18495   12.151 < 2.2e-16
SOUTH       4.85680     0.26892   18.061 < 2.2e-16

Lambda: 0.37676 LR test value: 241.96 p-value: < 2.22e-16
Asymptotic standard error: 0.024649 z-value: 15.285 p-value: < 2.22e-16
Wald statistic: 233.62 p-value: < 2.22e-16
```

```
Log likelihood: -9242.908 for error model
ML residual variance (sigma squared): 22.806, (sigma: 4.7756)
Number of observations: 3085
Number of parameters estimated: 4
AIC: 18494, (AIC for lm: 18734)
```

Here again, incorporating spatial effects does not substantially change the evidence. The dummy variable remains highly significant. Compared to the OLS results, its value is slightly lower, but the standard error is larger.

The fit of the two spatial models can be compared by means of the maximized log likelihood. For the error model this is `-9242.9` but for the lag model it is `-9232.9`, confirming its superiority. However, both models leave a lot unexplained and should only be used in an exploratory phase of possible spatial heterogeneity.

8.4   Spatial Regimes

Spatial regime regressions allow the model coefficients to vary between discrete spatial subsets of the data. We will implement this by taking advantage of the interaction feature of the formula in R (indicated by the : symbol between two sets of variables). Specifically, we will create a dummy variable for each regime (i.e., taking a value of one for observations in the regime and zero for all others) and then interact each explanatory variable with each dummy. A Chow test and its extension to spatial regressions (spatial Chow test) forms the basis for assessing the significance of the regimes.

8.4.1   Preliminaries

We create two new variables and add them to the `nat60` data frame. First, we construct the complement of the `SOUTH` dummy, `NOSOUTH`. Next, we construct a vector of ones of the same length as the other variables. This is easily done by adding `SOUTH` and `NOSOUTH` without having to worry about specifying the length. We call the new data frame `nat60a`.

```
> NOSOUTH = 1 - SOUTH
> ONE = SOUTH + NOSOUTH
> nat60a <- cbind(nat60,ONE,NOSOUTH)
```

Check that everything went fine using a `summary`. Note how the mean of `SOUTH` is 0.4577 and the mean of `NOSOUTH` is 0.5423. To have easy access to the variables, `attach(nat60a)` (detach other data sets if necessary).

8.4.2 Spatial Regimes in OLS

We will first run a base regression for the whole nation, in which the coefficients are kept fixed and without a `SOUTH` dummy. This will be referred to as the constrained model (since the coefficients are constrained to be equal across regimes). The corresponding output is:

```
> nat60ols <- lm(HR60 ~ RD60 + PS60 + MA60 +
+ DV60 + UE60)
> summary(nat60ols)

Call:
lm(formula = HR60 ~ RD60 + PS60 + MA60 + DV60 + UE60)

Residuals:
    Min      1Q  Median      3Q     Max
-14.999  -2.427  -0.821   1.434  90.954

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 10.51667    0.61064  17.222  < 2e-16 ***
RD60         2.69490    0.09653  27.919  < 2e-16 ***
PS60         0.60360    0.09141   6.603 4.72e-11 ***
MA60        -0.27212    0.01935 -14.063  < 2e-16 ***
DV60         1.29591    0.09599  13.501  < 2e-16 ***
UE60        -0.10610    0.03565  -2.976  0.00294 **
---
Signif. codes:  0 *** 0.001 ** 0.01 * 0.05 . 0.1   1

Residual standard error: 4.84 on 3079 degrees of freedom
Multiple R-Squared: 0.2673,Adjusted R-squared: 0.2661
F-statistic: 224.7 on 5 and 3079 DF,  p-value: < 2.2e-16
```

To obtain the unconstrained regression, we need to drop the default constant from the formula in order to make sure that a different intercept is estimated for each regime. This is done by specifying `0` in the formula for the regression specification. We also need to interact each of the explanatory

variables (including the constant, hence the need for `ONE`) with the `SOUTH` and `NOSOUTH` dummies.

We get the results from the familiar `lm` function with the customized formula, as shown:

```
> nat60rega <- lm(HR60 ~ 0 + (ONE + RD60 + PS60 +
+ MA60 + DV60 + UE60):(SOUTH + NOSOUTH))
> summary(nat60rega)

Call:
lm(formula = HR60 ~ 0 + (ONE + RD60 + PS60 + MA60 +
    DV60 + UE60):(SOUTH +  NOSOUTH))

Residuals:
     Min       1Q   Median       3Q      Max
-13.5265  -2.1820  -0.6467   1.3390  88.4821

Coefficients:
             Estimate Std. Error t value Pr(>|t|)
ONE:SOUTH     13.21547    0.87183  15.158  < 2e-16 ***
ONE:NOSOUTH    6.39963    0.99955   6.403 1.76e-10 ***
RD60:SOUTH     1.76448    0.15369  11.481  < 2e-16 ***
RD60:NOSOUTH   1.85730    0.23653   7.852 5.60e-15 ***
PS60:SOUTH     0.29930    0.16611   1.802  0.07166 .
PS60:NOSOUTH   0.37748    0.11730   3.218  0.00130 **
MA60:SOUTH    -0.27521    0.02951  -9.326  < 2e-16 ***
MA60:NOSOUTH  -0.19537    0.02978  -6.560 6.27e-11 ***
DV60:SOUTH     1.17945    0.18879   6.247 4.75e-10 ***
DV60:NOSOUTH   1.11883    0.11463   9.761  < 2e-16 ***
UE60:SOUTH    -0.29186    0.05518  -5.289 1.31e-07 ***
UE60:NOSOUTH   0.09277    0.04554   2.037  0.04173 *
---
Signif. codes:  0 *** 0.001 ** 0.01 * 0.05 . 0.1   1

Residual standard error: 4.721 on 3073 degrees of freedom
Multiple R-Squared: 0.5746,Adjusted R-squared: 0.5729
F-statistic: 345.9 on 12 and 3073 DF,  p-value: < 2.2e-16
```

Note that the `R-Squared` may not be computed correctly when the intercept is excluded in this fashion, but we can ignore that for now. The results

show some differences in significance between regimes for some of the coefficients, such as PS60 (not significant in SOUTH but significant in NOSOUTH) and UE60 (strongly significant and negative in SOUTH and positive but marginally significant in NOSOUTH).

8.4.3 Chow Test

A straightforward test on the constraint that the coefficients are equal across regimes is the Chow test contained in any standard econometrics test:

$$C = \frac{e_R' e_R - e_U' e_U}{k} \Big/ \frac{e_U' e_U}{N - 2k} \sim F(k, N - 2k), \tag{8.1}$$

where e_R is the vector of residuals from the constrained regression (constant coefficients) and e_U is the vector of residuals from the unconstrained regression (regime coefficients).

We extract the residuals from the regression object by means of the residuals method applied to the object.

```
> er <- residuals(nat60ols)
```

We also extract the degrees of freedom in a similar way.

```
> k <- nat60ols$rank
> n2k <- nat60ols$df.residual - k
```

The Chow test can be computed in a short function, shown below:

```
chow.test <- function(rest,unrest)
{
    er <- residuals(rest)
    eu <- residuals(unrest)
    er2 <- sum(er^2)
    eu2 <- sum(eu^2)
    k <- rest$rank
    n2k <- rest$df.residual - k
    c <- ((er2 - eu2)/k) / (eu2 / n2k)
    pc <- pf(c,k,n2k,lower.tail=FALSE)
    list(c,pc,k,n2k)
}
```

This function is contained in a file chow.test.R, which needs to be sourced to bring it into your workspace (make sure the file is in your working directory). Alternatively, you can type it in at the command line.

The function simply extracts the respective residual vectors from the regression objects, as well as the degrees of freedom and computes the simple statistic. Inference is based on the F distribution using the `pf` function. The results are passed back as a list containing the test statistic, p value and two degrees of freedom.

In our example, this gives:

```
> source("chow.test.R")
> chow.test(nat60ols,nat60rega)
[[1]]
[1] 27.11882
[[2]]
[1] 1.141437e-31
[[3]]
[1] 6
[[4]]
[1] 3073
```

The value of the test statistic is `27.12`, which is highly significant. Hence, there is strong evidence that the model coefficients indeed are not constant across regimes, indicating spatial heterogeneity.

8.4.4 Spatial Regimes in Spatial Regression

The residuals of the spatial regime regression continue to show strong evidence of spatial autocorrelation. Both `LMerr` and `LMlag` statistics are highly significant, as are their robust forms, with a slight edge in favor of the lag alternative.

```
> lm.LMtests(nat60rega,natq,
+ test=c("LMerr","RLMerr","LMlag","RLMlag"))


Lagrange multiplier diagnostics for spatial dependence

data:
model: lm(formula = HR60 ~ 0 + (ONE + RD60 + PS60 + MA60
 + DV60 + UE60):(SOUTH + NOSOUTH))
weights: natq
LMerr = 97.6755, df = 1, p-value < 2.2e-16


Lagrange multiplier diagnostics for spatial dependence
```

```
data:
model: lm(formula = HR60 ~ 0 + (ONE + RD60 + PS60 + MA60
 + DV60 + UE60):(SOUTH + NOSOUTH))
weights: natq
RLMerr = 9.8075, df = 1, p-value = 0.001738


Lagrange multiplier diagnostics for spatial dependence

data:
model: lm(formula = HR60 ~ 0 + (ONE + RD60 + PS60 + MA60
 + DV60 + UE60):(SOUTH + NOSOUTH))
weights: natq
LMlag = 129.1925, df = 1, p-value < 2.2e-16


Lagrange multiplier diagnostics for spatial dependence

data:
model: lm(formula = HR60 ~ 0 + (ONE + RD60 + PS60 + MA60
 + DV60 + UE60):(SOUTH + NOSOUTH))
weights: natq
RLMlag = 41.3245, df = 1, p-value = 1.289e-10
```

Spatial regimes can be incorporated in spatial lag and spatial error models in the same way as for OLS, by means of the interaction terms in the model formula.

For the spatial lag model, this gives:

```
> hr60rlag <- lagsarlm(HR60 ~ 0 + (ONE + RD60 + PS60 +
+ MA60 + DV60 + UE60):(SOUTH + NOSOUTH),data=nat60a,
+ natq)
> summary(hr60rlag)

Call:lagsarlm(formula = HR60 ~ 0 + (ONE + RD60 + PS60 +
    MA60 + DV60 +  UE60):(SOUTH + NOSOUTH),
    data = nat60a, listw = natq)

Residuals:
      Min        1Q    Median        3Q       Max
-13.41352  -2.11313  -0.62028   1.33659  88.38962
```

```
Type: lag
Coefficients: (asymptotic standard errors)
             Estimate Std. Error z value  Pr(>|z|)
ONE:SOUTH    9.739325   0.895945 10.8704  < 2.2e-16
ONE:NOSOUTH  5.273614   0.982158  5.3694 7.899e-08
RD60:SOUTH   1.420458   0.153772  9.2374  < 2.2e-16
RD60:NOSOUTH 1.625286   0.231632  7.0167 2.272e-12
PS60:SOUTH   0.208484   0.162155  1.2857  0.198544
PS60:NOSOUTH 0.385196   0.114367  3.3681  0.000757
MA60:SOUTH  -0.213140   0.029087 -7.3277 2.340e-13
MA60:NOSOUTH -0.165157  0.029204 -5.6553 1.556e-08
DV60:SOUTH   0.979970   0.184224  5.3195 1.041e-07
DV60:NOSOUTH 0.898503   0.112754  7.9687 1.554e-15
UE60:SOUTH  -0.210014   0.053885 -3.8974 9.722e-05
UE60:NOSOUTH 0.083777   0.044455  1.8845  0.059494

Rho: 0.25961 LR test value: 106.05 p-value: < 2.22e-16
Asymptotic standard error: 0.02568 z-value: 10.11 p-value: < 2.22e-16
Wald statistic: 102.21 p-value: < 2.22e-16

Log likelihood: -9106.562 for lag model
ML residual variance (sigma squared): 21.189, (sigma: 4.6032)
Number of observations: 3085
Number of parameters estimated: 14
AIC: 18241, (AIC for lm: 18345)
LM test for residual autocorrelation
test value: 38.109 p-value: 6.6914e-10
```

Here again, there is evidence of a difference between regimes for the coefficients of PS60 and UE60. In order to assess this more rigorously, we also need the results for the spatial lag model of the constrained specification. The corresponding result is:

```
> hr60lag <- lagsarlm(HR60 ~ RD60 + PS60 + MA60 +
+ DV60 + UE60,data=nat60a,natq)
> summary(hr60lag)

Call:lagsarlm(formula = HR60 ~ RD60 + PS60 + MA60 + DV60 + UE60,
    data = nat60a,     listw = natq)

Residuals:
```

```
         Min         1Q     Median         3Q        Max
-13.89718   -2.17000   -0.74723    1.37653   90.09303


Type: lag
Coefficients: (asymptotic standard errors)
            Estimate Std. Error z value  Pr(>|z|)
(Intercept)  7.048613   0.629963 11.1889 < 2.2e-16
RD60         1.912544   0.106597 17.9419 < 2.2e-16
PS60         0.456590   0.088178  5.1781 2.242e-07
MA60        -0.190804   0.019269 -9.9023 < 2.2e-16
DV60         0.942877   0.093493 10.0850 < 2.2e-16
UE60        -0.052096   0.034187 -1.5239    0.1275


Rho: 0.32721 LR test value: 192.26 p-value: < 2.22e-16
Asymptotic standard error: 0.024022 z-value: 13.621 p-value: < 2.22e-16
Wald statistic: 185.53 p-value: < 2.22e-16


Log likelihood: -9143.04 for lag model
ML residual variance (sigma squared): 21.528, (sigma: 4.6399)
Number of observations: 3085
Number of parameters estimated: 8
AIC: 18302, (AIC for lm: 18492)
LM test for residual autocorrelation
test value: 50.95 p-value: 9.4769e-13
```

8.4.5 Spatial Chow Test

When the regression error terms are spatially autocorrelated, the simple form of the Chow test shown in (8.1) is no longer valid. A corrected version of the test is referred to as a *spatial* Chow test (see Anselin 1988, pp. 123–124). With the results of both the constrained and unconstrained models at hand, this can be implemented as a LR test, since the constrained model is nested within the unconstrained one.

This can be readily implemented in a small function:

```
spatialchow.test <- function(rest,unrest)
{
    lrest <- rest$LL
    lunrest <- unrest$LL
    k <- rest$parameters - 2
```

```
        spchow <- - 2.0 * (lrest - lunrest)
        pchow <- pchisq(spchow,k,lower.tail=FALSE)
        list(spchow,pchow,k)
}
```

The function takes the spatial regression objects and extracts the maximized log-likelihood as the `$LL` item. The degrees of freedom are obtained as the number of parameters in the constrained model less two (for the autoregressive parameter and the error variance), using `$parameters`. The computation of the LR test is straightforward. Significance is assessed from the Chi-squared distribution. The results are returned as a list.

This function is contained in the file `spatialchow.test.R`. Source this file to incorporate it into your workspace and invoke the test by passing the two spatial lag regression objects to it: `hr60lag` for the restricted model and `hr60rlag` for the unrestricted model. The result shows a strong rejection of the null hypothesis, suggesting significantly different coefficients in each of the regimes, even after correcting for the spatial lag.

```
> source("spatialchow.test.R")
> spatialchow.test(hr60lag,hr60rlag)
[[1]]
          [,1]
[1,] 72.95666
[[2]]
              [,1]
[1,] 1.010420e-13
[[3]]
[1] 6
```

The approach for a spatial error model is similar, by passing the appropriate constrained and unconstrained models to the `spatialchow.test` function.

8.5   Practice

To practice, assess the spatial heterogeneity for other years of the US homicide data set, using a spatial Chow test for both lag and error specifications.

Exercise 9

Continuous Spatial Heterogeneity

9.1 Objectives

We continue the analysis of spatial heterogeneity in two forms: spatial expansion and geographically weighted regression (GWR). These both yield individual coefficient estimates for each location. This exercise illustrates how you can manipulate the regression formula to implement spatial expansion and introduces the basic GWR approach, highlighting the sensitivity of the results to the choice of bandwidth and kernel function. We will map the location-specific coefficient estimates in *GeoDa* and *ArcMap*.

9.2 Preliminaries

For the spatial expansion example, we will use the built-in `columbus` data set to keep matters simple. Make sure you have the `spdep` library loaded. We won't actually use its functionality, but we need it for easy access to the Columbus sample data.

For GWR, we will use the package `spgwr`:[1]

```
> library(spgwr)
```

---

[1]There is also a specialized software package `GWR 3.x`, distributed by Professor Stewart Fotheringham at the National Centre for Geocomputation of the National University of Ireland. It can be obtained from http://ncg.nuim.ie/ncg/GWR/software.htm. For any in-depth analysis using GWR, the use of this more comprehensive software package is recommended. The `spgwr` package currently only replicates a small subset of its functionality.

If you want to make maps of the individual coefficient estimates, make sure you have *GeoDa* or *ArcMap* handy, with the `columbus.shp` shape file.

9.3 Spatial Expansion Method

As the first example of modeling continuous spatial heterogeneity, we consider Casetti's expansion method. This is a special case of a model with varying coefficients, where the variation is driven by an *expansion equation*.

In the simplest example, the expansion equation is a linear trend surface. This yields a *final equation* that contains the original explanatory variables, as well as interaction terms with the X and Y coordinates. A straightforward calculation yields individual coefficient values for each location, which can then be mapped and further analyzed.

A common problem in the implementation of the spatial expansion method is the high degree of multicollinearity that results from the interaction terms. In practice, this requires the use of some adjustments, such as replacing the original interaction terms by their principal components. Such adjustments are outside the current scope and we will ignore the potential multicollinearity.

We will use the built-in `columbus` data set to illustrate this procedure.

Start by loading the data set and attaching it.

```
> data(columbus)
> attach(columbus)
```

We will run the standard regression example, with `CRIME` as the dependent variable and `INC` and `HOVAL` as the explanatory variables. The expansion will consist of a linear trend in the coordinates `X` and `Y`.

To create the proper interaction terms, we use the `cross` form of the formula (symbolized by `*`). The final regression is then as shown.

```
> colex <- lm(CRIME ~ (INC + HOVAL)*(X + Y))
> summary(colex)

Call:
lm(formula = CRIME ~ (INC + HOVAL) * (X + Y))

Residuals:
     Min       1Q   Median       3Q      Max
-18.5556  -7.6351  -0.6181   7.8363  30.1948
```

```
Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 108.97559   69.36676   1.571   0.1241
INC          -5.82949    3.84408  -1.516   0.1373
HOVAL         0.27337    0.82049   0.333   0.7407
X            -0.76287    1.13692  -0.671   0.5061
Y            -0.26332    1.21420  -0.217   0.8294
INC:X        -0.01854    0.05396  -0.344   0.7329
INC:Y         0.13949    0.08004   1.743   0.0891 .
HOVAL:X       0.03159    0.01549   2.040   0.0480 *
HOVAL:Y      -0.05034    0.02196  -2.293   0.0272 *
---
Signif. codes:  0 *** 0.001 ** 0.01 * 0.05 . 0.1   1

Residual standard error: 11.04 on 40 degrees of freedom
Multiple R-Squared: 0.6375,Adjusted R-squared: 0.5649
F-statistic: 8.791 on 8 and 40 DF,  p-value: 7.663e-07
```

Note how very few coefficients are shown to be significant, even though the R-squared is more than acceptable. This is a typical symptom of high multicollinearity.

For our purposes, we will ignore this potential problem and compute individual coefficients for HOVAL for each location. This is accomplished by extracting the relevant coefficients from the regression object using the $coefficients item.

```
> b <- colex$coefficients
> b[3]
     HOVAL
0.2733714
> b[8]
   HOVAL:X
0.03159375
> b[9]
    HOVAL:Y
-0.05034417
```

The individual coefficients are obtained as;

```
> bihoval <- b[3] + b[8] * X + b[9] * Y
> bihoval
```
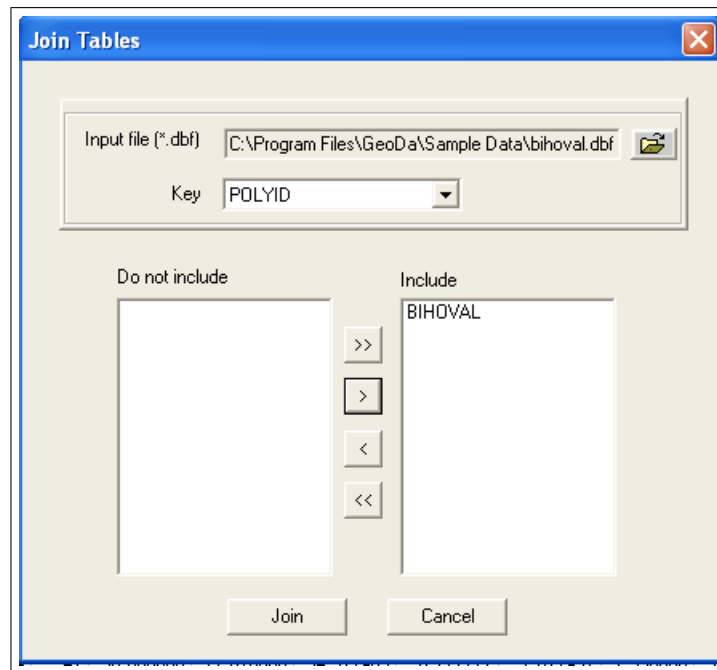
Figure 9.1: Join tables dialog in *GeoDa*

```
 [1] -0.71945869 -0.73484522 -0.54173838
 . . .
 [49]  0.36488625
> summary(bihoval)
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
-0.7348 -0.3418 -0.1479 -0.1096  0.1362  0.6105
```

We see that the varying coefficients range from `-0.7348` to `0.6105` (compare this to an estimate of `-0.2739` in the fixed coefficient model).

We can now create a data frame that can be written out to a dbf file, to be joined with the dbf file of the columbus shape file for mapping.

```
> bi <- data.frame(POLYID,bihoval)
> write.dbf(bi,"bihoval.dbf")
```

In *GeoDa*, load the `columbus.shp` shape file and select the `table`. Right click and choose `join table`. In the dialog, enter the name of the dbf file `bihoval.dbf`, the Key as POLYID, and the variable to be joined as BIHOVAL, as illustrated in Figure 9.1 .

Figure 9.2: Map of spatially expanded coefficients for `HOVAL`

Now, the new variable is available to all the analyses and visualization methods in *GeoDa*. For example, we can create a simple quantile map using 6 quantiles to show the gradual change of the coefficients across the city neighborhoods, with the low values in the upper left, and the highest values in the lower right, as in Figure 9.2. Note that the spatial expansion does not need to be limited to the use of `X` and `Y` coordinates, but can be applied to any set of variables that may underlie the spatial heterogeneity in the model.

9.4 GWR

9.4.1 Basics

Make sure the `spgwr` library is active and the `columbus` data set `attached`. We first explore the basics of the `gwr` function. This function creates an object that contains a data frame with the individual coefficient estimates for each variable and each location. The arguments to this function are a

formula (for the model specification), a data set, a matrix with the X, Y coordinates of the locations (or polygon centroids) and a bandwidth. The default kernel function is the Gaussian.

The bandwidth is hard to specify a priori and the preferred approach is to carry out a cross-validation, minimizing the root mean square prediction error. For now, we'll try three values of 20, 3 and 2. As the bandwidth widens, the surface of estimated coefficients becomes smoother.

With the familiar CRIME regression, the coordinates as the X and Y variables from the built in data set and the bandwidth of 20, this yields:

```
> colg1 <- gwr(CRIME ~ INC + HOVAL,data=columbus,
+ coords=cbind(X,Y),bandwidth=20)
> colg1
Call:
gwr(formula = CRIME ~ INC + HOVAL, data = columbus,
    coords = cbind(X, Y), bandwidth = 20)
Kernel function: gwr.gauss
Fixed bandwidth: 20
Summary of GWR coefficient estimates:
                Min. 1st Qu.  Median 3rd Qu.    Max.  Global
X.Intercept. 68.0700 69.0300 69.4500 69.8000 71.1700 68.6190
INC          -1.7200 -1.6590 -1.6210 -1.5600 -1.4510 -1.5973
HOVAL        -0.3407 -0.3078 -0.2789 -0.2454 -0.1903 -0.2739
```

The print command gives a summary of the estimated coefficients (note that summary is not useful in this context), listing the basic parameters (kernel and bandwidth) and providing descriptive statistics for each of the estimated coefficients. For example, the estimates for HOVAL range from -0.3407 to -0.1903 and are not that far from the global estimate of -0.2739. The main reason for this is the large bandwidth, which does not allow for much variability. Compare this to the range obtained for the spatial expansion method of -0.7348 to 0.6105.

For the two other bandwidths, we get the following results:

```
> colg2 <- gwr(CRIME ~ INC + HOVAL,data=columbus,
+ coords=cbind(X,Y),bandwidth=3)
> colg2
Call:
gwr(formula = CRIME ~ INC + HOVAL, data = columbus,
    coords = cbind(X, Y), bandwidth = 3)
Kernel function: gwr.gauss
```

```
Fixed bandwidth: 3
Summary of GWR coefficient estimates:
                 Min.  1st Qu.   Median  3rd Qu.     Max.  Global
X.Intercept. 23.03000 53.33000 62.96000 68.72000 81.39000 68.6190
INC          -3.22200 -1.88300 -0.75950 -0.34360  1.30700 -1.5973
HOVAL        -1.07400 -0.38630 -0.12820  0.04305  0.85320 -0.2739
> colg3 <- gwr(CRIME ~ INC + HOVAL, data=columbus,
+ coords=cbind(X,Y),bandwidth=2)
> colg3
Call:
gwr(formula = CRIME ~ INC + HOVAL, data = columbus,
    coords = cbind(X, Y), bandwidth = 2)
Kernel function: gwr.gauss
Fixed bandwidth: 2
Summary of GWR coefficient estimates:
                 Min.  1st Qu.   Median  3rd Qu.     Max.  Global
X.Intercept. 20.87000 43.06000 65.07000 71.34000 88.53000 68.6190
INC          -6.01900 -1.74400 -0.71910 -0.03815  3.85400 -1.5973
HOVAL        -3.47700 -0.46420 -0.20710  0.04728  1.40400 -0.2739
```

The bandwidth of 3 yields a similar range to the expansion method, while the bandwidth of 2 results in a very wide range. The three distributions can be compared to that of the expansion method by means of a box plot.

To extract the vector of coefficients we need to access the SDF attribute of the gwr class. This is a data frame containing the coordinates, coefficients and a measure of error. For example, for the second bandwidth (2), this looks like (the R2 and gwr.e columns have been dropped):

```
> colg3$SDF
        coordinates    sum.w X.Intercept.         INC        HOVAL ...
1   (38.8, 44.07) 1.146372      46.18933 -0.70880055 -0.207112745 ...
2  (35.62, 42.38) 1.425847      43.06483 -0.66509539 -0.222184163 ...
3  (39.82, 41.18) 1.269489      53.52575 -0.97462747 -0.248571427 ...
4   (36.5, 40.52) 1.918875      58.23655 -0.22800553 -0.678264390 ...
...
49 (42.67, 24.96) 1.883877      74.88421  3.85392142 -3.477123284 ...
```

A vector with values for an individual coefficient is extracted using the familiar $ notation for a list item, e.g.:

```
> hovg3 <- colg2$SDF$HOVAL
```

Figure 9.3: Box plots for different bandwidths in GWR

```
> summary(hovg3)
    Min.  1st Qu.   Median     Mean  3rd Qu.     Max.
-1.07400 -0.38630 -0.12820 -0.14260  0.04305  0.85320
```

We create vectors for the other two bandwidths in a similar way:

```
> hovg20 <- colg1$SDF$HOVAL
> summary(hovg20)
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
-0.3407 -0.3078 -0.2789 -0.2768 -0.2454 -0.1903
> hovg2 <- colg3$SDF$HOVAL
```

```
> summary(hovg2)
    Min.  1st Qu.   Median     Mean  3rd Qu.     Max.
-3.47700 -0.46420 -0.20710 -0.25020  0.04728  1.40400
```

We can now put box plots for the three coefficient vectors side by side.

```
> boxplot(bihoval,hovg20,hovg3,hovg2,
+ names=c("Expansion","bw=20","bw=3","bw=2"))
```

This yields Figure 9.3. Note how the distribution for bandwidth 20 is concentrated around the median value (a very smooth surface), whereas the distribution for bandwidth 2 shows several outliers. The distribution for the linear expansion methods is similar in range to that for a bandwidth of 3, but the median is lower.

Just as for the expansion method coefficients, we can group the vectors for GWR into a data frame, export it to a `dbf` file to join with the `columbus.shp` shape file.

```
> colgwr <- data.frame(POLYID,hovg20,hovg3,hovg2)
> write.dbf(colgwr,"colgwr.dbf")
```

We can then map these, say in *ArcMap*, as illustrated in Figures 9.4 and 9.5.[2] Here, we show a choropleth map using natural breaks with 8 categories. This clearly illustrates the smoothing effect of the larger bandwidth.

9.4.2   Selecting the Optimal Bandwidth

So far, the bandwidth we specified for the `gwr` procedure was purely ad hoc. **spgwr** implements two methods to select an *optimal* bandwidth. The most commonly used is cross validation, minimizing the mean squared prediction error when each observation is left out in turn. A second method is based on the minimization of the Akaike Information Criterion (AIC) (for technical details, see Fotheringham et al. 2002, pp. 59–62).

The optimal bandwidth is returned by the `gwr.sel` function, for a given formula, data set, coordinate matrix and kernel (the default is the Gaussian kernel, `gwr.gauss`). The default method is cross-validation (`method="cv"`). In order to use the AIC criterion, the method argument must be set explicitly, as `method="aic"`.

---

[2] *GeoDa* is less appropriate for this purpose, since in the current version the join function only allows one additional variable to be added. To add more variables, use `Excel` to join and save as `dbf` before loading the shape file into *GeoDa*. Alternatively, one could use the mapping capabilities built into the `sp` package when the data set is a spatial data frame.
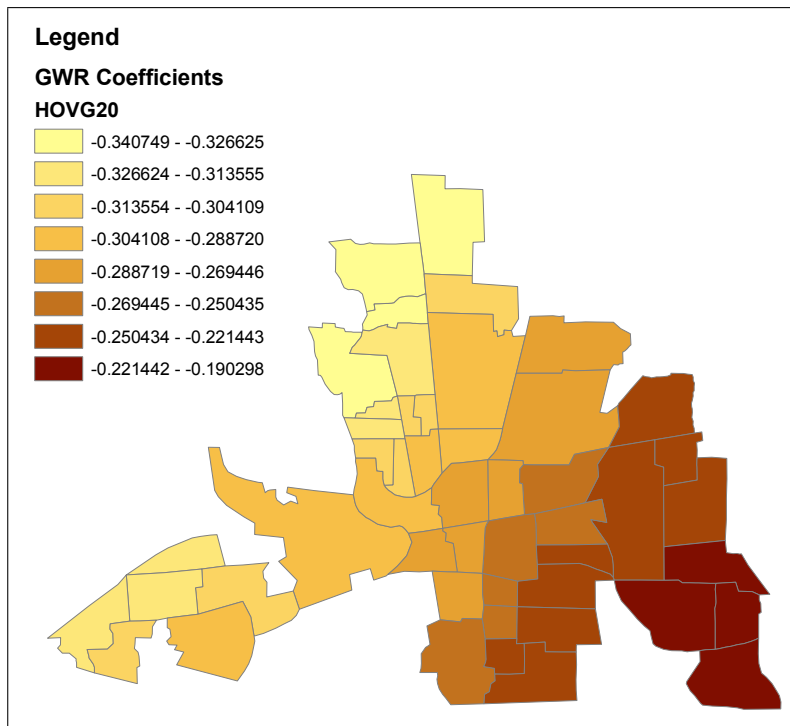
Figure 9.4: GWR coefficients for `HOVAL`, bandwidth=20

This suggests a two step procedure to implement GWR: first, run `gwr.sel` to obtain the bandwidth, next use this value in a final run of `gwr`. In our example, with cross validation first, this yields:

```
> bw1 <- gwr.sel(CRIME ~ INC + HOVAL,data=columbus,
+ coords=cbind(X,Y))
Bandwidth: 12.65221 CV score: 12.25277
Bandwidth: 20.45127 CV score: 12.32858
Bandwidth: 7.83213 CV score: 12.22244
Bandwidth: 4.853154 CV score: 11.7606
Bandwidth: 3.012046 CV score: 11.15676
Bandwidth: 1.874179 CV score: 14.61981
Bandwidth: 3.715287 CV score: 11.24019
Bandwidth: 3.041249 CV score: 11.14678
Bandwidth: 3.276828 CV score: 11.12378
Bandwidth: 3.444304 CV score: 11.15169
```
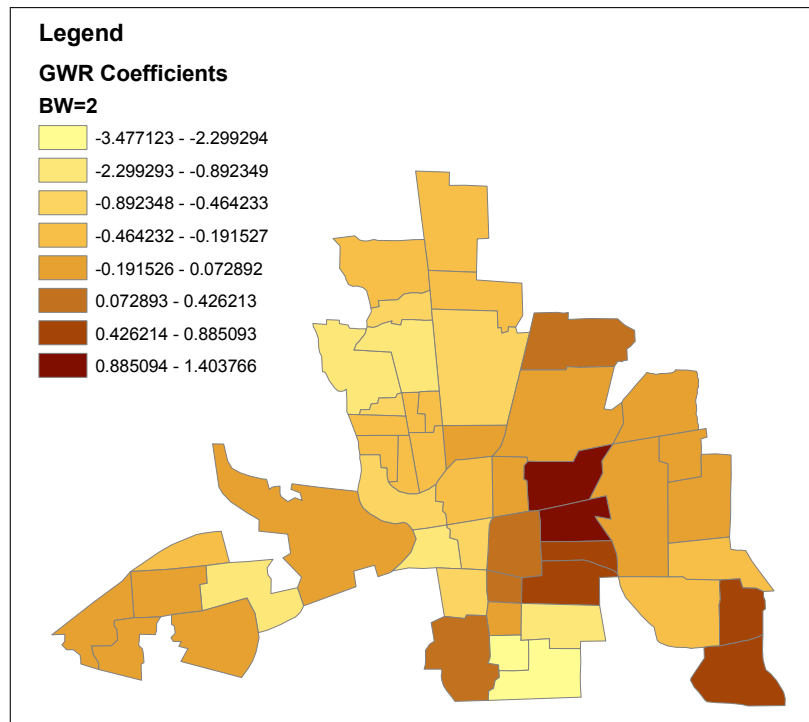
Figure 9.5: GWR coefficients for `HOVAL`, bandwidth=2

```
Bandwidth: 3.233495 CV score: 11.12159
Bandwidth: 3.22236 CV score: 11.12143
Bandwidth: 3.216871 CV score: 11.12141
Bandwidth: 3.217402 CV score: 11.12141
Bandwidth: 3.217443 CV score: 11.12141
Bandwidth: 3.217484 CV score: 11.12141
Bandwidth: 3.217443 CV score: 11.12141
> bw1
[1] 3.217443
```

The second step follows as:

```
> colg4 <- gwr(CRIME ~ INC + HOVAL,data=columbus,
+ coords=cbind(X,Y),bandwidth=bw1)
> colg4
Call:
gwr(formula = CRIME ~ INC + HOVAL, data = columbus,
```

```
      coords = cbind(X, Y), bandwidth = bw1)
Kernel function: gwr.gauss
Fixed bandwidth: 3.217443
Summary of GWR coefficient estimates:
                 Min.   1st Qu.   Median  3rd Qu.     Max.   Global
X.Intercept. 23.23000 54.12000 63.90000 68.76000 80.90000 68.6190
INC          -3.13100 -1.91300 -0.98440 -0.36860  1.29100 -1.5973
HOVAL        -1.05300 -0.37670 -0.09739  0.03004  0.79460 -0.2739
```

To use the AIC criterion, we include the `method` argument explicitly in the call to `gwr.sel`:

```
> bw2 <- gwr.sel(CRIME ~ INC + HOVAL,data=columbus,
+ coords=cbind(X,Y),method="aic")
Bandwidth: 12.65221 AIC: 382.8966
Bandwidth: 20.45127 AIC: 383.3532
Bandwidth: 7.83213 AIC: 382.1
Bandwidth: 4.853154 AIC: 381.4758
Bandwidth: 3.012046 AIC: 412.4025
Bandwidth: 5.991021 AIC: 380.7834
Bandwidth: 6.106824 AIC: 380.8622
Bandwidth: 5.718077 AIC: 380.6517
Bandwidth: 5.387706 AIC: 380.6667
Bandwidth: 5.578766 AIC: 380.6282
Bandwidth: 5.573107 AIC: 380.6281
Bandwidth: 5.563802 AIC: 380.628
Bandwidth: 5.564886 AIC: 380.628
Bandwidth: 5.564926 AIC: 380.628
Bandwidth: 5.564967 AIC: 380.628
Bandwidth: 5.564926 AIC: 380.628
> colg4 <- gwr(CRIME ~ INC + HOVAL,data=columbus,
+ coords=cbind(X,Y),bandwidth=bw2)
> colg4
Call:
gwr(formula = CRIME ~ INC + HOVAL, data = columbus,
    coords = cbind(X, Y), bandwidth = bw2)
Kernel function: gwr.gauss
Fixed bandwidth: 5.564926
Summary of GWR coefficient estimates:
                Min. 1st Qu.  Median 3rd Qu.    Max.  Global
X.Intercept. 30.0400 62.2700 69.5000 70.9200 75.3900 68.6190
```

```
INC            -2.5810 -2.0010 -1.6170 -0.8033 -0.2798 -1.5973
HOVAL          -0.8032 -0.3798 -0.1042 -0.0458  0.2542 -0.2739
```

Note how the AIC criterion results in a much smoother parameter map
than the one obtained with cross validation. There is little evidence as to
which is *superior* since they optimize different objective functions. Since our
use of GWR is primarily as a diagnostic for spatial heterogeneity, a careful
sensitivity analysis is recommended in any case.

As an exercise, compare the distribution of the HOVAL (or other) coeffi-
cient between the two optimal bandwidths and to its counterpart obtained
with the spatial expansion method in a box plot and choropleth map.

9.4.3   Selecting a Kernel Function

A second important practical aspect of the use of GWR is the selection of the
kernel function that is at the basis of the computation of the local estimates.
This is carried out by means of the gweight argument to the gwr.sel and
gwr functions. So far, we have used the Gaussian kernel, but spgwr also
implements a bisquare kernel. In contrast to the Gaussian kernel, which has
an infinite range (the bandwidth is not an absolute cut off), the bisquare
kernel sets all values beyond the bandwidth to zero. While the choice of a
kernel function should not be arbitrary, in the literature the emphasis is on
the sensitivity of the results to the choice of a bandwidth, rather than the
kernel function.

With the cross-validation approach an optimal bandwidth of 33 is se-
lected, which is much larger than for the Gaussian kernel. The resulting
GWR estimates for HOVAL are roughly comparable to the ones obtained
with a bandwidth of 20 for the Gaussian kernel.

```
> bwbs1 <- gwr.sel(CRIME ~ INC + HOVAL,data=columbus,
+ coords=cbind(X,Y),gweight=gwr.bisquare)
Bandwidth: 12.65221 CV score: 12.92097
Bandwidth: 20.45127 CV score: 12.41531
Bandwidth: 25.27136 CV score: 12.37858
Bandwidth: 23.70165 CV score: 12.38804
Bandwidth: 28.25033 CV score: 12.36507
Bandwidth: 30.09144 CV score: 12.36080
Bandwidth: 31.69725 CV score: 12.35831
Bandwidth: 31.08388 CV score: 12.35917
Bandwidth: 32.22175 CV score: 12.35763
Bandwidth: 32.54591 CV score: 12.35724
```

```
Bandwidth: 32.74625 CV score: 12.35701
Bandwidth: 32.87007 CV score: 12.35687
Bandwidth: 32.94660 CV score: 12.35678
Bandwidth: 32.99389 CV score: 12.35673
Bandwidth: 33.02312 CV score: 12.35670
Bandwidth: 33.04119 CV score: 12.35668
Bandwidth: 33.05235 CV score: 12.35666
Bandwidth: 33.05925 CV score: 12.35666
Bandwidth: 33.06351 CV score: 12.35665
Bandwidth: 33.06615 CV score: 12.35665
Bandwidth: 33.06778 CV score: 12.35665
Bandwidth: 33.06879 CV score: 12.35664
Bandwidth: 33.06941 CV score: 12.35664
Bandwidth: 33.06979 CV score: 12.35664
Bandwidth: 33.07003 CV score: 12.35664
Bandwidth: 33.07018 CV score: 12.35664
Bandwidth: 33.07027 CV score: 12.35664
Bandwidth: 33.07032 CV score: 12.35664
Bandwidth: 33.07037 CV score: 12.35664
Bandwidth: 33.07037 CV score: 12.35664
> colg5 <- gwr(CRIME ~ INC + HOVAL,data=columbus,
+ coords=cbind(X,Y),gweight=gwr.bisquare,
+ bandwidth=bwbs1)
> colg5
Call:
gwr(formula = CRIME ~ INC + HOVAL, data = columbus,
    coords = cbind(X, Y), bandwidth = bwbs1,
    gweight = gwr.bisquare)
Kernel function: gwr.bisquare
Fixed bandwidth: 33.07037
Summary of GWR coefficient estimates:
                Min. 1st Qu.  Median 3rd Qu.    Max.  Global
X.Intercept. 68.3100 68.9600 69.2600 69.5700 71.6400 68.6190
INC          -1.7270 -1.6480 -1.6140 -1.5690 -1.4680 -1.5973
HOVAL        -0.3331 -0.3052 -0.2806 -0.2535 -0.1883 -0.2739
```

The AIC criterion yields a much smaller optimal bandwidth of 11.

```
> bwbs2 <- gwr.sel(CRIME ~ INC + HOVAL,data=columbus,
+ coords=cbind(X,Y),gweight=gwr.bisquare,method="aic")
```

```
Bandwidth: 12.65221 AIC: 382.7242
Bandwidth: 20.45127 AIC: 384.3786
Bandwidth: 7.83213 AIC: 386.7498
Bandwidth: 15.27372 AIC: 384.0654
Bandwidth: 10.81111 AIC: 381.6533
Bandwidth: 9.673238 AIC: 383.0491
Bandwidth: 11.25258 AIC: 381.6384
Bandwidth: 11.07023 AIC: 381.6049
Bandwidth: 11.05193 AIC: 381.6044
Bandwidth: 11.04548 AIC: 381.6044
Bandwidth: 11.04647 AIC: 381.6044
Bandwidth: 11.04651 AIC: 381.6044
Bandwidth: 11.04655 AIC: 381.6044
Bandwidth: 11.04651 AIC: 381.6044
> colg6 <- gwr(CRIME ~ INC + HOVAL,data=columbus,
+ coords=cbind(X,Y),gweight=gwr.bisquare,
+ bandwidth=bwbs2)
> colg6
Call:
gwr(formula = CRIME ~ INC + HOVAL, data = columbus,
    coords = cbind(X, Y), bandwidth = bwbs2,
    gweight = gwr.bisquare)
Kernel function: gwr.bisquare
Fixed bandwidth: 11.04651
Summary of GWR coefficient estimates:
                Min.  1st Qu.  Median  3rd Qu.    Max.  Global
X.Intercept. 27.04000 63.63000 69.37000 71.06000 77.33000 68.6190
INC          -2.57400 -2.03700 -1.65300 -0.66770 -0.26290 -1.5973
HOVAL        -0.73720 -0.43290 -0.08757 -0.04483  0.26540 -0.2739
```

As before, the individual parameter estimates can be extracted from the
`gwr` object, plotted or mapped.

9.5 Practice

To practice, plot or map the coefficient vector for the other coefficients in
the model. Alternatively, check for continuous spatial heterogeneity in the
BOSTON or BALTIMORE data sets. Compare the insights provided by the
expansion method to those from GWR, and carry out sensitivity analysis
for the choice of bandwidth and kernel function.

Appendix

Summary of Data and Spatial Weights Operations

- Data import and setup

  - `library(spdep)`: load the `spdep` package library
  - `source("read.geoda.R")`: load and compile the source file *read.geoda.R*
  - `read.geoda`: read a data set exported from *GeoDa*
  - `data(dataset)`: load an R data set
  - `attach(dataset)`: attach the data frame, making the variables available by name
  - `summary(dataset)`: summary descriptive statistics of all variables in the data set

- Types of spatial weights objects

  - `nb`: neighbor list, contiguities only
  - `listw`: neighbor list with weights, sparse format
  - `knn`: k-nearest neighbor object
  - `matrix`: standard matrix object

- Reading spatial weights files into neighbor lists

  - `read.gal`: read GAL format weights files
  - `read.gwt2nb`: read GWT format weights files

- Creating spatial weights objects

  - `cell2nb(nrow,ncol)`: neighbor list for rectangular regular grid cells

- **dnearneigh(coordinatematrix,lowdist,updist)**: create distance band weights using the `lowdist` and `updist` limits
- **knearneigh(coordinatematrix,k)**: create a matrix with the indices of the k-nearest neighbors
- **nblag(nb,maxlag)**: higher order neighbors
- **poly2nb(polygonlist)**: neighbor object from polygon boundary file
- **tri2nb(coordinatematrix)**: create neighbor list using Thiessen polygon neighbor structure

- Conversion between weights objects

  - **knn2nb(knn)**: create a neighbor list object from a `knn` object (a matrix)
  - **nb2listw(nb)**: create a `listw` weights object from a neighbor list
  - **mat2listw(matrix)**: convert a spatial weights n by n matrix to a `listw` object
  - **nb2mat(nb)**: convert a neighbor list to a matrix
  - **make.sym.nb(nb)**: force a possibly asymmetric neighbor list (e.g., from k-nearest neighbors) to a symmetric form

- Characteristics of spatial weights

  - **summary(nb)**: connectivity summary of neighbor list
  - **card(nb)**: number of neighbors for each observation
  - **is.symmetric.nb(nb)**: check on symmerty of connectivity structure

- Spatial transformations

  - **lag.listw(listw,variables)**: create spatially lagged variables using a `listw` object
  - **invIrM(listw,rho)**: create inverse matrix $(I - \rho W)^{-1}$

132

Bibliography

Anselin, L. (1988). *Spatial Econometrics: Methods and Models.* Kluwer Academic Publishers, Dordrecht, The Netherlands.

Anselin, L. (2003a). *Data and Spatial Weights in spdep, Notes and Illustrations.* http://sal.agecon.uiuc.edu/stuff_main.php#tutorials.

Anselin, L. (2003b). *An Introduction to Spatial Regression Analysis in R.* http://sal.agecon.uiuc.edu/stuff_main.php#tutorials.

Anselin, L. (2005). *Exploring Spatial Data with GeoDa: A Workbook.* Spatial Analysis Laboratory (SAL). Department of Geography, University of Illinois, Urbana-Champaign, IL.

Baller, R., Anselin, L., Messner, S., Deane, G., and Hawkins, D. (2001). Structural covariates of U.S. county homicide rates: Incorporating spatial effects. *Criminology*, 39(3):561–590.

Bivand, R. S. (2001). More on spatial data analysis. *R News*, 1(3):13–17.

Bivand, R. S. (2002a). Implementing spatial data analysis software tools in R. In Anselin, L. and Rey, S., editors, *New Tools for Spatial Data Analysis: Proceedings of the Specialist Meeting.* Center for Spatially Integrated Social Science (CSISS), University of California, Santa Barbara. CD-ROM.

Bivand, R. S. (2002b). Spatial econometrics functions in R: Classes and methods. *Journal of Geographical Systems*, 4:405–421.

Bivand, R. S. and Gebhardt, A. (2000). Implementing functions for spatial statistical analysis using the R language. *Journal of Geographical Systems*, 2:307–317.

Bivand, R. S. and Portnov, B. A. (2004). Exploring spatial data analysis techniques using R: The case of observations with no neighbors. In Anselin, L., Florax, R. J., and Rey, S. J., editors, *Advances in Spatial Econometrics: Methodology, Tools and Applications*, pages 121–142. Springer-Verlag, Berlin.

Fotheringham, A. S., Brunsdon, C., and Charlton, M. (2002). *Geographically Weighted Regression*. John Wiley, Chichester.

Kelejian, H. H. and Robinson, D. P. (1992). Spatial autocorrelation: A new computationally simple test with an application to per capita country police expenditures. *Regional Science and Urban Economics*, 22:317–333.

Ord, J. K. (1975). Estimation methods for models of spatial interaction. *Journal of the American Statistical Association*, 70:120–126.

Venables, W., Smith, D., and R Development Core Team (2004). *An Introduction to R – Notes on R, a Programming Environment for Data Analysis and Graphics, Version 2.0.1 (2004-11-15)*. http://cran.r-project.org.